



DEPARTMENT OF INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

Benchmarking Supersingular Isogeny Diffie-Hellman Implementations

Jonas Hagg





DEPARTMENT OF INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

Benchmarking Supersingular Isogeny Diffie-Hellman Implementations

Benchmarking Supersingular Isogeny Diffie-Hellman Implementierungen

Author:	Jonas Hagg
Supervisor:	Prof. Dr. Claudia Eckert
Advisor:	Prof. Dr. Daniel Loebenberger
Submission Date:	Submission date



I confirm that this bachelor's thesis in informatics is my own work and I have documented all sources and material used.

Munich, Submission date

Jonas Hagg

Acknowledgments

Abstract

Kurzfassung

Contents

Acknowledgments	iii
Abstract	iv
Kurzfassung	v
1. Background	1
1.1. Key Exchange	3
1.1.1. Diffie-Hellman Key-Exchange	3
1.1.2. Key Encapsulation	4
1.1.3. Differences	5
1.2. Post-Quantum Cryptography	5
1.2.1. Impact of Quantum Computers on Cryptography	6
1.2.2. Classes of Post-Quantum Cryptography	8
1.2.3. Isogeny-based Cryptography	9
2. Description of existing SIDH implementations	16
2.1. SIKE	16
2.2. PQCrypto-SIDH	17
2.3. CIRCL	17
3. Benchmarks	19
3.1. Benchmarking Methodology	19
3.2. Application Details	21
3.2.1. Application Flow	22
3.2.2. Application Structure	23
3.2.3. Adding Implementations	26
3.3. Usage	27
3.4. Results	29
3.4.1. SIKE Implementations	30
3.4.2. Optimized Implementations	31
3.4.3. Compressed implementations	33
4. Performance Analysis	35
4.1. Measures for Efficiency	35
4.1.1. Comparing SIDH libraries	35
4.1.2. Comparing SIDH and ECDH	37

4.2. Security Considerations	37
4.2.1. Constant time	37
4.2.2. Key length	37
5. Conclusion	38
6. Introduction	39
6.1. Section	39
6.1.1. Subsection	39
A. General Addenda	42
A.1. Detailed Benchmarks	42
A.1.1. Benchmarks for ECDH	42
A.1.2. Benchmarks for Sike Reference	43
A.1.3. Benchmarks for Sike Generic	44
A.1.4. Benchmarks for Sike Generic Compressed	45
A.1.5. Benchmarks for Sike x64	46
A.1.6. Benchmarks for Sike x64 Compressed	47
A.1.7. Benchmarks for CIRCL x64	48
A.1.8. Benchmarks for Microsoft Generic	49
A.1.9. Benchmarks for Microsoft Generic Compressed	50
A.1.10. Benchmarks for Microsoft x64	51
A.1.11. Benchmarks for Microsoft x64 Compressed	52
List of Figures	53
List of Tables	54
Bibliography	55

1. Background

In modern cryptography one can differ between symmetric and asymmetric encryption schemes. While in a symmetric scheme the decryption and encryption of data is processed with the same key, asymmetric protocols introduce a key pair for every participant: A public key for encryption and a private key for decryption. The public key of asymmetric protocols is, as the name suggests, public to everyone. However, the private key needs to be secret and nobody but the producer may have knowledge about the private key.

Symmetric Cryptography

In Figure 1.1 a classical symmetric encryption scheme is shown. First, a plaintext is encryption using a symmetric encryption algorithm and a secret key. The resulting chiffré text is transported to the receiver, where it is decrypted using the appropriate decryption algorithm and the same secret key. The red section in the middle represents an insecure channel (e.g. the internet), where attackers may read or modify data. A critical question is the exchange of the key through the insecure channel. Somehow the symmetric key needs to be transported securely to the receiver of the chiffré.

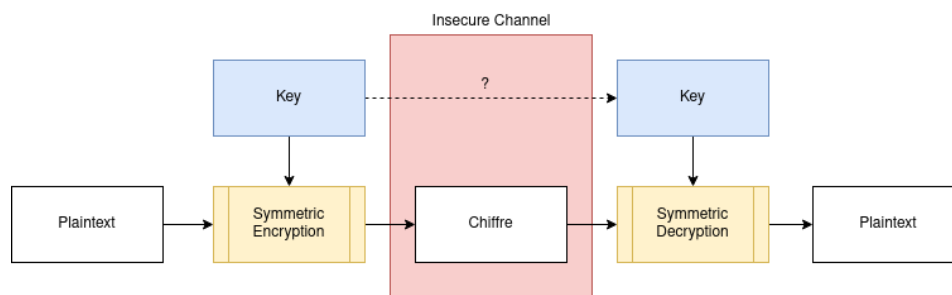


Figure 1.1.: Simple symmetric encryption scheme: Encryption and decryption algorithm use the same key.

In the following, the symmetric decryption and encryption is expressed in a more formal way. The common key k is used for encryption (Enc) and decryption (Dec), p is the plaintext and c is the ciphertext:

$$Enc(p, k) = c$$

$$Dec(c, k) = p$$

Asymmetric Cryptography

In asymmetric cryptography each participating subject needs to generate a key pair, that consist of a private key and a public key. As mentioned above, the public key needs to be public (e.g. stored in a public database or a public key server). The private key is only known to the subject and needs to be kept secret. Figure 1.2 shows an example for asymmetric encryption. Assume Alice wants to send encrypted data to Bob. Therefore, Bob created a key pair and published his public key. Alice requests Bobs public key (e.g. from a public database) and uses it to encrypt the data. Once Bob received the ciphertext, he uses his secret private key for decryption in order to retrieve the original plaintext. In this work, the term *public-key encryption* or *public-key algorithm* is used as a synonym for asymmetric cryptography.

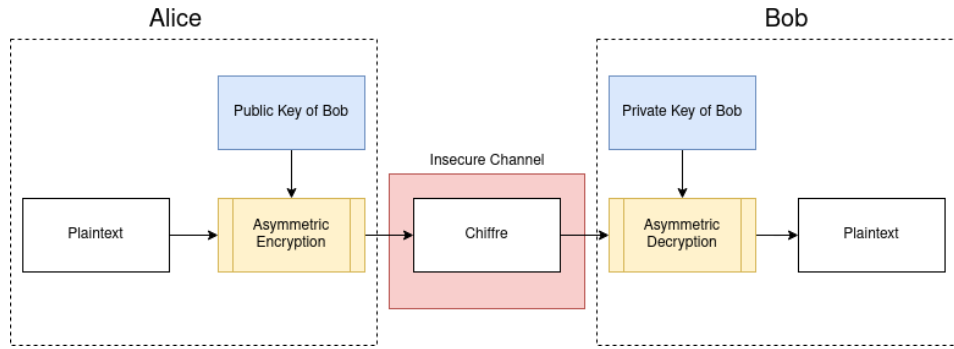


Figure 1.2.: Asymmetric encryption scheme: Encryption and decryption algorithm use different keys.

To formalize this procedure, again assume p as plaintext and c as ciphertext. The generated key pair of Bob consists of a private key for decryption (d_{Bob}) and a public key for encryption (e_{Bob}).

$$\begin{aligned} Enc(p, e_{Bob}) &= c \\ Dec(c, d_{Bob}) &= p \end{aligned}$$

In contrast to the symmetric encryption, no secret key needs to be exchanged. However, the encryption and decryption of data using asymmetric encryption require intensive mathematical computations. Hence, the encryption of big sets of data using asymmetric encryption is not efficient. On the other hand, symmetric encryption algorithms are usually based on simple operations, such as bit shifting or XOR. This can be implemented efficiently in software and hardware. Thus, the practical relevance of symmetric encryption is enormous [1].

As stated above, securely exchanged keys are a precondition for the use of efficient symmetric encryption schemes. In order to exchange arbitrary keys securely, different key exchange protocols are available.

1.1. Key Exchange

1.1.1. Diffie-Hellman Key-Exchange

The Diffie-Hellman key exchange was introduced by Whitfield Diffie and Martin Hellman in 1976 [2]. The resulting shared key of the protocol is calculated decentralized and is never transported over an insecure channel.

Protocol

The classical Diffie-Hellman key exchange assumes, that Alice and Bobs want to create a shared secret key. Therefore, they agree on a big prime p and g , which is a primitive root modulo p ¹. Both, p and g are not secret and may be known to the public [3].

1. Alice choses a random $a \in \{1, 2, \dots, p - 2\}$ as private key.
2. Alice calculates the public key $A = g^a \bmod p$.
3. Bob choses a random $b \in \{1, 2, \dots, p - 2\}$ as private key.
4. Bob calculates the public key $B = g^b \bmod p$.
5. Alice and Bob exchange their public keys A and B .
6. Alice calculates:

$$\begin{aligned} k_{AB} &= B^a \bmod p \\ &= (g^b \bmod p)^a \bmod p \\ &= g^{ab} \bmod p \end{aligned} \tag{1.1}$$

7. Bob calculates:

$$\begin{aligned} k_{AB} &= A^b \bmod p \\ &= (g^a \bmod p)^b \bmod p \\ &= g^{ba} \bmod p \\ &= g^{ab} \bmod p \end{aligned} \tag{1.2}$$

8. Alice and Bob created the shared secret k_{AB} . Note that only the public keys of Alice and Bob were send over an insecure channel. The generated secret was calculated decentralized by Alice and Bob.

This procedure also can be illustrated in the following diagram, which emphasizes the commutative properties of the protocol. It does not matter if first apply $x \rightarrow x^a$ or $x \rightarrow x^b$ to the starting point g . The result is the same, since $g^{ab} = g^{ba}$.

¹The primitive root modulo p is a generator element for the set $S = \{1, 2, \dots, p - 1\}$ [1].

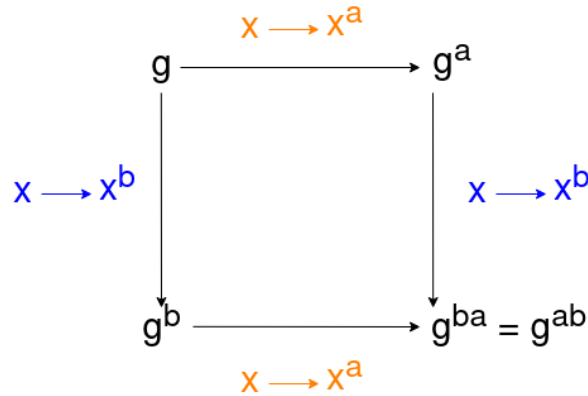


Figure 1.3.: Diffie Hellman diagram, both paths lead to the same result.

Security

If an attacker wants to compute k_{AB} , it needs to compute the private keys A and B of Alice and Bob. Since only the public keys are exchanged, the attacker needs to compute:

$$b = \log_g B \bmod p$$

$$a = \log_g A \bmod p$$

Hence, the security of the classical Diffie-Hellman key exchange is based on the discrete logarithm problem (see 1.2). When using ephemeral key pairs, the Diffie-Hellman key exchange may be used as efficient perfect forward secrecy (PFS) protocol [1].

In modern cryptography elliptic curves are often used to increase the security of the Diffie-Hellman key exchange (ECDH). The participants have to agree on an elliptic curve and a point P on the elliptic curve. The computations to generate the shared secret k_{AB} follow the same principles, but obviously they are adopted to work on elliptic curves. The advantage of ECDH is the increased security strength while using the same key length than the classical Diffie-Hellman protocol [1].

Note that the introduced protocol does not authorize the participating subjects and does not guarantee integrity. Thus, this simple protocol may be exploited by a Man-In-The-Middle attack. A more advanced protocol using certificates and signed messages could be implemented, to guarantee authentication and integrity [1].

1.1.2. Key Encapsulation

Key encapsulation methods (KEM) transmit a previously generated symmetric key. KEMs usually use asymmetric key pairs in order to encrypt the generated symmetric key. In the following the concept is shown by the RSA based PKCS #1 v1.5 algorithm, which uses RSA key pairs to transmit a shared secret from Alice to Bob [4]:

1. Bob generates a RSA key pair with the public key e_{Bob} and the private key d_{Bob} and transmits the public key to Alice.
2. Alice generates a random secret key k_{AB} :

$$k_{AB} = \text{random}()$$

3. Alice maps this secret to an integer m , using a well-defined mapping function h :

$$m = h(k_{AB})$$

4. Alice encrypts m with Bobs public key using the RSA encryption algorithm and transmits c to Bob.

$$c = \text{RSA}_{enc}(m, e_{Bob})$$

5. Bob decrypts the received ciphertext s to obtain the integer m :

$$m = \text{RSA}_{dec}(c, d_{Bob})$$

6. Finally bob uses the inverse mapping function h^{-1} to retrieve the shared secret:

$$k_{AB} = h^{-1}(m)$$

1.1.3. Differences

TODO: differences between kem and kex (decentralized, pfs?), discrete log vs factorization

1.2. Post-Quantum Cryptography

In the following a *classical computer* refers to a non-quantum computer, which can be simulated by a deterministic Turing machine. In contrast to *classical computer* the term *quantum computer* describes a machine, which uses quantum mechanical phenomena to perform computations. It is important to note, that quantum computers can simulate classical computers. In addition, classical computer are able to simulate quantum computers with exponential time overhead. Thus, classical and quantum computations can calculate the same class of functions. However, quantum computers enable operations, that allow much faster computation [5].

In the past, scientists queried, if large-scale quantum computer are a physical possibility. It was stated, that the underlying quantum states are too fragile and hard to control. [6] Today, quantum error correction codes are known, which put a large-scale quantum computers within the realms of possibility [7]. However, it is still a big engineering challenge from a laboratory approach to a general-purpose quantum computer, which involves thousands or millions of physical qubits [6].

The security of modern asymmetric cryptographic primitives is usually based on difficult

number theoretic problems, e.g. the discrete logarithm problem (DH, ECDH) or the factorization problem (RSA) [6]. These problems are theoretically solvable, but the computation of a result on classical computer claims an impractical amount of resources. In 2019, scientists solved the factorization problem for a 240 digit integer in about 900 core-years on a classical computer (one core year means running a CPU for a full year) [8]. In the following the discrete logarithm problem and the factorization problem are described.

Discrete Logarithm Problem

The discrete logarithm problem is the following challenge [9]: Given a prim p and two integers g and y . Find an integer x , such that

$$y = g^x \bmod p \\ \iff x = \log_g y \bmod p$$

Until today, it is not known if a classical computer is able to compute the general discrete logarithm problem in polynomial time. Thus, the discrete logarithm problem is considered to be difficult so solve for classical computers [9]. This assumption makes the discrete logarithm problem an attractive basis for DSA, ElGamal, the classical Diffie-Hellman protocol, and for the elliptic curve Diffie-Hellman (ECDH).

Factorization Problem

Given two large primes p and q , it is easy to compute the the product of them:

$$n = p \star q$$

For a given n , however, it is difficult to find the prime factors p and q . The computation of the prime factorization for a given integer n is called the factorization problem [1]. For large numbers n no efficient algorithm for classical computers is known to solve this challenge [1]. The most famous cryptographic protocol, which builds upon the hardness of the factorization problem is RSA.

1.2.1. Impact of Quantum Computers on Cryptography

As stated above, quantum computers enable new operations which speed up certain algorithms. Two quantum algorithms which have enormous consequences on modern cryptography are *Shor's algorithm* and *Grover's algorithm* [5].

Shor's Algorithm

Peter Shor published "*Algorithms for quantum computation: discrete logarithms and factoring*" in 1994 [10], where he demonstrated that the factorization problem and the discrete logarithm problem can be solved in polynomial time on quantum computers. Both problems are the basis of many public-key systems (RSA, DH, ECDH, ...), which are used intensively in modern

communication systems. Hence, a quantum computer running *Shor's algorithm* would qualify the assumption of most asymmetric encryption schemes and thus break their security.

Grover's Algorithm

The second algorithm having impact on computer security was published by Lov Grover in 1996 ("*A fast quantum mechanical algorithm for database search*", [11]) - namely *Grover's algorithm*. The algorithm solves the problem of finding an element y in a set s (e.g. a database) where $|s| = N$. On a classical computer an algorithm solving the problem runs in $\mathcal{O}(N)$, however *Grover's algorithm* has complexity $\mathcal{O}(\sqrt{N})$ [5].

In contrast to public-key systems, which rely on hard mathematical problems, symmetric encryption schemes rely on the secrecy of a randomly generated key. Thus, to break symmetric encryption one needs to perform a brute-force attack on the symmetric key. Using *Grover's algorithm* offers a square root speed up on classical brute force attacks [12]. Assume a randomly generated n -bit key. A classical brute force algorithm lies in $\mathcal{O}(2^n)$, which is considered to be safe for a big n (e.g. $n=128$). *Grover's algorithm* speeds up this attack to $\mathcal{O}(\sqrt{2^n}) = \mathcal{O}(2^{n/2})$ [12]. However, the complexity is still exponential and with a growing key size n the security can be further increased. Thus, *Grover's algorithm* forces symmetric encryption schemes to increase their key size in order to stay secure.

To sum up, quantum computers make use of quantum mechanical phenomena in order to solve mathematical problems, which are assumed to be difficult for modern computers. As a result large-scale quantum computers might break many algorithms of modern asymmetric cryptography and enforce increased key sizes for symmetric encryption schemes. The following table from the NIST "*Report on Post-Quantum Cryptography*" [6] shows the impact of quantum computers on modern encryption schemes:

Cryptographic algorithm	Type	Purpose	Impact from quantum computer
AES	Symmetric key	Encryption	Larger key sizes needed
SHA	—	Hash functions	Larger output needed
RSA	Public key	Signatures, key establishment	No longer secure
DSA	Public key	Signatures, key exchange	No longer secure
ECDH, ECDSA	Public key	Signatures, key exchange	No longer secure

Table 1.1.: Impact of quantum computers on modern encryption schemes

In contrast to this development in modern cryptography NIST states [6]:

"In the last three decades, public key cryptography has become an indispensable component

of our global communication digital infrastructure. These networks support a plethora of applications that are important to our economy, our security, and our way of life, such as mobile phones, internet commerce, social networks, and cloud computing. In such a connected world, the ability of individuals, businesses and governments to communicate securely is of the utmost importance."

This statement emphasizes the urgency and need of new asymmetric encryption schemes. As a consequence, NIST initiated a process to standardize quantum-secure public-key algorithms. In literature, this is called *post-quantum cryptography*, since the objectives of the submitted procedures is to stay secure against a large-scale quantum computer. In July 2020, Round 3 of the standardization process was announced. Different approaches for quantum-resistant algorithms have been proposed. In this work, the focus is on isogeny-based cryptography (section 1.2.3). Other classes of post-quantum cryptography are described for completeness in the following section 1.2.2.

1.2.2. Classes of Post-Quantum Cryptography

This section provides an overview about important post-quantum cryptography classes: Lattice-based, multivariate and code-based cryptography as well as hash-based signatures are shortly presented.

Lattice-based Cryptography

Lattice-based cryptography is - as the name suggests - based on the mathematical construct of lattices². There are different computational optimization problems involving lattices, which are considered to be hard to solve even for quantum computers [13]. In 1998, NTRU was published as the first public-key system based on lattices [14]. Since then NTRU was continuously improved resulting in NTRUencrypt (public-key system) and NTRUsign (digital signing algorithm). Furthermore, a fully homomorphic encryption scheme based on lattices was published in 2009 [15].

Lattice-based cryptography is characterized by simplicity and efficiency [6]. The security of existing implementations (NTRU or Ring-LWE) can be reduced to NP-hard problems. However, lattices encryption schemes have problems to prove security against known cryptanalysis [6].

Multivariate Cryptography

Multivariate cryptography are public key systems that are based on multivariate polynomials (e.g. $p(x, y) = x + 2y$) over a finite field \mathbb{F} . Their security is based on the prove, that solving systems of multivariate polynomials are NP-hard [16]. This makes multivariate public key systems attractive for post-quantum cryptography; especially their short signatures make

²A lattice l is a subgroup of \mathbb{R}^n . In the context of cryptography usually integer lattices are considered: $l \subseteq \mathbb{Z}^n$ [13].

them a candidate for quantum-secure digital signature algorithms [17], e.g. the Rainbow signature scheme [18].

Code-based Cryptography

Code-based cryptographic primitives are build upon error-correcting codes. A public key system using error-correcting codes uses a public key to add errors to a given plaintext resulting in a ciphertext. Only the owner of the private key is able to correct there errors and to reconstruct the plaintext [19]. McEliece, published in 1978, was the first of those systems and it has not been broken until today [20]. Beside asymmetric cryptography, code-based schemes have been proposed for digital signatures, random number generators and cryptographic hash functions [19].

Hash-based Signatures

Hash-based signatures describe the construction of digital signatures schemes based on hash functions. Thus, the security of theses primitives is based on the security of the underlying hash function and not on hard algorithmic problems [19]. Since hash functions are widely deployed in modern computer systems, the security of hash-based signatures is well understood [6].

The initially developed One-Time Signatures have the downside, that a new public key pair is needed for each signature [21]. In 1979, Merkle introduced the Merkle Signature Scheme (MSS), which uses one public key for multiple signatures [22]. Further improvements of MSS introduced public keys, which can be used for 2^{80} signatures. However, this also leads to longer signature sizes [21].

1.2.3. Isogeny-based Cryptography

Isogeny-based cryptography was proposed in 2011 as a new cryptographic system, that might resists quantum computing [23]. Beside the publication describing isogeny-based cryptography, the authors also provided a reference implementaion of a public key system, which is called SIKE [24]. Isogeny-based cryptography benfits from small key sizes compared to other post-quantum cryptography classes, however, their performance is comparatively slow [24]. The security of these primitives is based on finding isogenies between supersingular elliptic curves.

In the following, the problem is roughly illustrated. It is not indented to precise the mathematics behind isogeny-based cryptography, since this is not the scope of this work. However, the reader might become a little understanding of the magic behind supersingular isogenies. Next, the central component of isogeny-based cryptography, namely the Supersingular Isogeny Diffie Hellman (SIDH), is described. Finally, details of the reference implementation SIKE will be given.

Illustration of the Problem [25] [26]

Isogeny-based cryptography works on supersingular elliptic curves. To be more precise, it is based on isogenies between supersingular elliptic curves.

In the context of elliptic curves, one can calculate a quotient of an elliptic curve E by a subgroup S . This essentially means to construct a new elliptic curve E/S . Beside this new curve, the procedure also yields a function $\phi_S : E \rightarrow E/S$, which is called *isogeny*. Carefully chosen elliptic curves E have a wide range of subgroups, which can be used to construct many isogenies.

Supersingular elliptic curves are a special type of elliptic curves, which have properties that are useful for cryptography. Since supersingular elliptic curves can be seen as subset from ordinary elliptic curves, they can also be used to calculate isogenies between certain curves, as described above.

The idea behind isogeny-based cryptography might be illustrated as followed:

1. Start with a known curve E and build a isogeny to a arbitrary reachable curve E_A .
2. This yields the isogeny $\phi_A : E \rightarrow E_A$, which is used as private key.
3. The curve E_A is used as part of the public key.

As usually in asymmetric cryptography, the hard mathematical problem is the calculation of the private key while knowing the public key. To be more precise: Find the isogeny $\phi_A : E \rightarrow E_A$ while knowing curves E and E_A is considered to be a quantum-resistant challenge. In literature, this is formally defined as *SIDH problem* [24].

This key pair, however, is not used to decrypt or encrypt data. In fact, the procedure is very similar to the previously introduced Diffie-Hellman key exchange, where the key pair is used to establish a shared secret between two communication partners. In its core, isogeny-based cryptography creates a shared secret via a Diffie-Hellman like procedure. This is called *Supersingular Isogeny Diffie Hellman (SIDH)*.

Supersingular Isogeny Diffie Hellman (SIDH)

Recall the Diffie Hellman diagram in Figure 1.3, where the commutative property of the protocol was shown. The same diagram can be drawn for the Diffie Hellman on supersingular isogenies, see Figure 1.4.

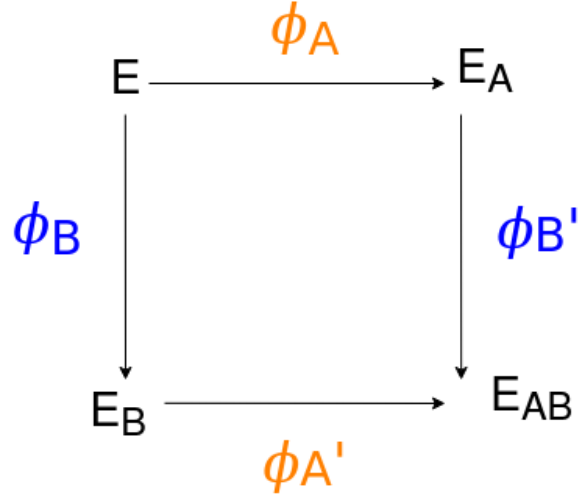


Figure 1.4.: Supersingular Isogeny Diffie Hellman procedure, both paths lead to the same result. Note, the different isogenies ϕ'_A and ϕ'_B , that are applied in each second step of the diagram. The reason for this lies in the mathematics of supersingular isogenies: $\phi_A\phi_B \neq \phi_B\phi_A$. In order to construct ϕ'_X or ϕ'_B , the public key of each communication partner provides additional information beside the curve E_A or E_B . [27]

In the previous section the underlying mathematical idea was illustrated. The described key generation can be extended to a key exchange primitive, which has strong similarity to the classical Diffie-Hellman key exchange. Starting point of SIDH is a known supersingular elliptic curve E .

1. Alice creates an isogeny ϕ_A (private key), which leads to curve E_A (part of the public key).
2. Bob creates an isogeny ϕ_B (private key), which leads to curve E_B (part of the public key).
3. Alice and Bob exchange their public keys.
4. Bob computes ϕ'_B (using additional information from the public key of Alice) and applies ϕ'_B to the received E_A . This results in E_{AB}
5. Alice computes ϕ'_A (using additional information from the public key of Bob) and applies ϕ'_A to the received E_B . This results in E_{AB}
6. Alice and Bob now share the common secret E_{AB} .

Implementation Details

The reference implementation of SIKE [24] provides two fundamental functions: *isogen* and *isoex*. Both are used, to implement the previously introduced Supersingular Isogeny Diffie-Hellman (SIDH) algorithm. Note, that the secret key sk is represented as a random integer.

Function	Input	Output
<i>isogen</i>	secret key sk	public key pk
<i>isoex</i>	secret key sk , public key pk	shared secret sec

Table 1.2.: Core functions of the SIKE reference implementation.

The function *isogen* takes a secret key (random integer) as input and generates the public key. The shared secret is generated by *isoex*, which takes the own secret key and the foreign public key as input. The key exchange procedure with respect to *isogen* and *isoex* works as followed:

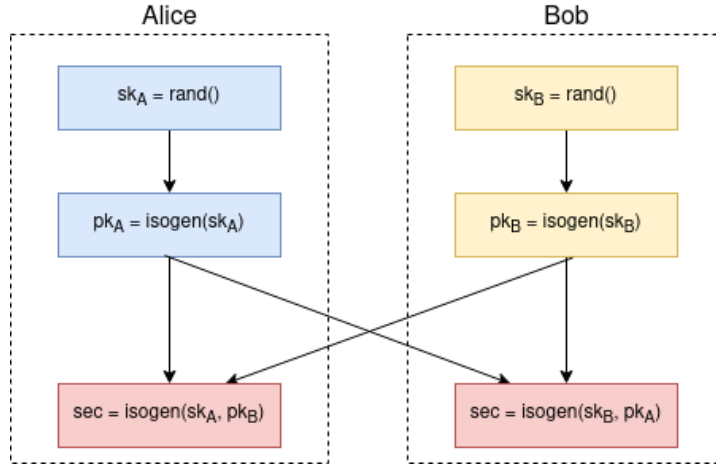


Figure 1.5.: SIDH based on *isogen* and *isoex*: After generating a random secret key sk , each party computes its public key pk . After the exchange of public keys, each party finally calculates the shared secret sec .

Beside this key exchange algorithm, SIKE provides a complete asymmetric encryption scheme and a key encapsulation mechanism (KEM) [24]. Both of these schemes build upon the here described SIKE core functions *isogen* and *isoex*.

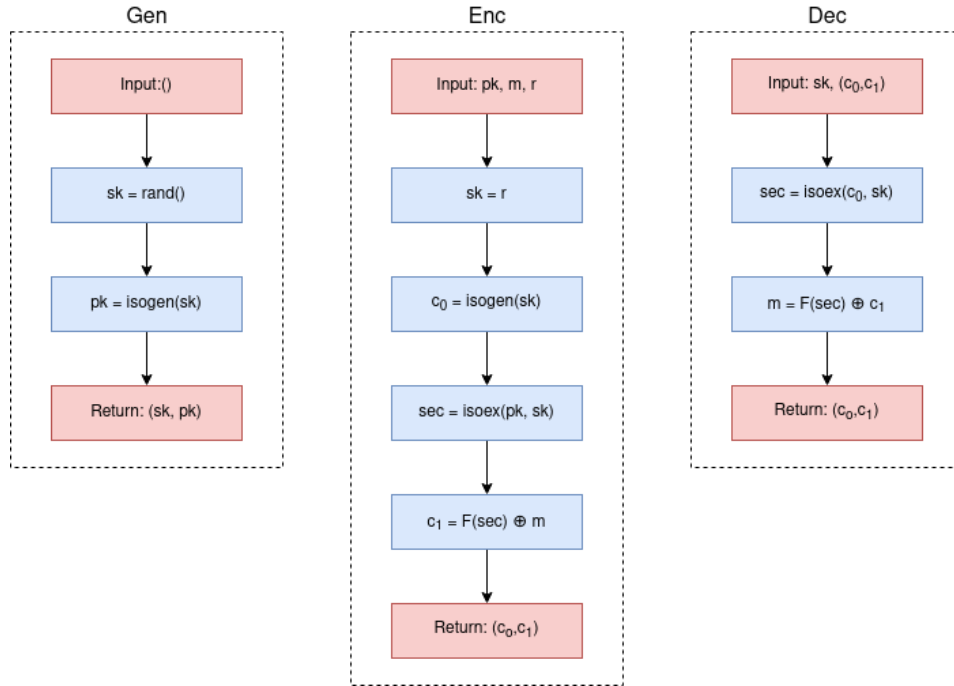


Figure 1.6.: The SIKE public key encryption system consists of three algorithms:

1. *Gen* generates a key pair (sk, pk) .
 2. *Enc* encrypts a given plaintext m using a foreign public key pk and the own secret key r .
 3. *Dec* decodes a given ciphertext using the own secret key sk .
- Note, that the function F used in *Enc* and *Dec* is a key derivation function.

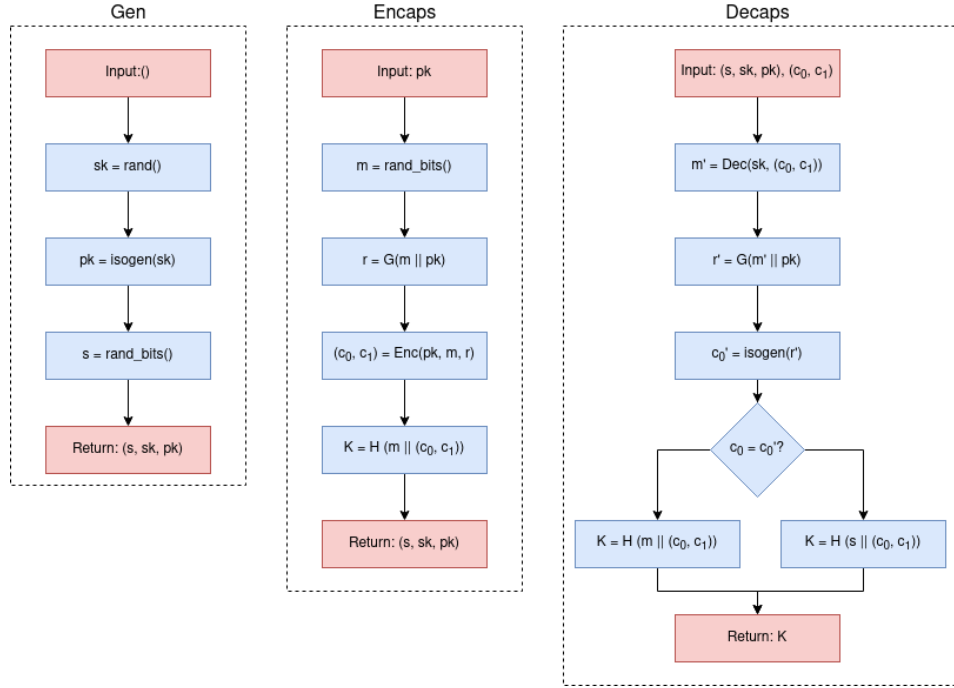


Figure 1.7.: The SIKE key exchange mechanism consists of three algorithms:

1. *Gen* TODO
2. *Encaps* TODO
3. *Decaps* TODO

Security considerations

The security of SIKE is based on the hardness of the *SIDH problem*: Given two supersingular elliptic curves E and E' , find an isogeny between them [24].

The SIKE reference implementation proposes different parameter sets, each supposing to ensure a NIST defined security level. These NIST defined security levels are:

1. Any attack breaking this security level must require resources comparable to perform a key search on a 128-bit key (e.g. AES128).
2. Any attack breaking this security level must require resources comparable to perform a collision search search on a 256-bit hash function (e.g. SHA256).
3. Any attack breaking this security level must require resources comparable to perform a key search on a 192-bit key (e.g. AES192).
4. Any attack breaking this security level must require resources comparable to perform a collision search search on a 384-bit hash function (e.g. SHA384).
5. Any attack breaking this security level must require resources comparable to perform a key search on a 256-bit key (e.g. AES256).

The proposed parameter sets of SIKE are named after the bit length of the underlying primes:

- SIKEp434 supposed to satisfy NIST security level 1 (AES128)
- SIKEp503 supposed to satisfy NIST security level 2 (SHA256)
- SIKEp610 supposed to satisfy NIST security level 3 (AES192)
- SIKEp751 supposed to satisfy NIST security level 5 (AES256)

Current research provides confidence, that the SIKE parameter sets satisfy the defined security levels even under the assumption of currently known algorithms [28]. Therefore, the authors consider three algorithms to solve the *SIDH problem*: *Tani's quantum claw finding algorithm* [29], *Grover's algorithm* [11] and a *parallel collision-finding algorithm* [30].

Ephemeral SIDH keys are predestined to implement quantum-secure perfect forward secrecy protocols (PFS) [longa2018note]. PFS ensures, that a compromised long-term key does not reveal past or future keys of the protocol.

Side-channel attacks against isogeny-based cryptography might 1) reveal parts of the secret private key or 2) reveal parts of the public key computation. To protect against power-analysis side-channel attacks it is recommended to prefer constant-time implementations [24]. The authors state, however, that an attacker has "*access to a wide range of power, timing, fault and various other side-channels*". Thus, preventing isogeny-based cryptography from all side-channel attacks seems to be nearly impossible.

2. Description of existing SIDH implementations

Currently, three implementations of Supersingular Isogeny Diffie-Hellman (SIDH) are available, namely *SIKE*, *CIRCL* and *PQCrypto-SIDH*. In this chapter each implementation is introduced in detail. At the end of the chapter, Table 2.1 shows similarities and differences between all approaches.

In the following, some algorithms are described as *compressed*. These compressed version exploit shorter public key sizes while increasing the computation time of the algorithms.

2.1. SIKE

SIKE stands for Supersingular Isogeny Key Encapsulation. It is the reference implementation of the first proposed isogeny-based cryptographic primitives [23]. Today, SIKE is a NIST candidate for quantum-resist "*Public-key Encryption and Key-establishment Algorithms*". It is developed by a cooperation of researchers, lead by David Jao [24].

SIKE implements its key encapsulation mechanism (KEM) upon a public key encryption system (PKE), which is built upon SIDH (as highlighted in chapter 1). Beside a generic reference implementation, SIKE offers various optimized implementations of their cryptographic primitives:

- Generic optimized implementation, written in portable C
- x64 optimized implementation, partly written in x64 assembly
- x64 optimized compressed implementation, partly written in x64 assembly
- ARM64 optimized implementation, partly written in ARMv8 assembly
- ARM Cortex M4 optimized implementation, partly written in ARM thumb assembly
- VHDL implementation

All of these implementations can be run with the following parameter sets: p434, p503, p610 and p751. SIKE states to countermeasure timing and cache attacks by implementing constant time cryptography [24].

2.2. PQCrypto-SIDH

PQCrypto-SIDH is a software library mainly written in C. It is developed by Microsoft for experimental purposes [31]. Note, that many developers of SIKE also work for Microsoft leading to great similarities between SIKE and PQCrypto-SIDH. However, in terms of compression, SIKE references the here described Microsoft library in its documentation. The PQCrypto-SIDH library implements a isogney-based KEM and the underlying SIDH. Moreover, the library offers the following optimized versions:

- Generic optimized implementation, written in portable C
- Generic optimized compressed implementation, written in portable C
- x64 optimized implementation, partly written in assembly
- x64 optimized compressed implementation, partly written in assembly
- ARMv8 optimized implementation, partly written in assembly
- ARMv8 optimized compressed implementation, partly written in assembly

All of these implementations can be run with the following parameter sets: p434, p503, p610 and p751. The developers argue to protect the algorithms against timing and cache attacks. Therefore, the library implements constant time operations on secret key material [31].

2.3. CIRCL

CIRCL (Cloudflare Interoperable, Reusable Cryptographic Library) is a by Cloudflare developed collection of cryptographic primitives [32]. CIRCL is written in Go and implements some quantum-secure algorithms like SIDH and an isogeny-based KEM. Cloudflare does not guarantee for any security within their library. Furthermore, the isogeny-based cryptographic primitives are adopted from the official SIKE implementation. The following implementation optimizations are stated to be available:

- Generic optimized implementation, written in Go (unfortunately, this version could not be compiled)
- AMD64 optimized implementation, partly written in assembly
- ARM64 optimized implementation, partly written in assembly

Note, that there are no compressed versions available. The library supports the following parameter sets: p434, p503 and p751. To avoid side-channel attacks, their code is implemented in constant time [33].

2. Description of existing SIDH implementations

	SIKE	PQCrypto-SIDH	CIRCL
Developer	Research cooperation	Microsoft	Cloudflare
Language	C Assembly	C Assembly	GO Assembly
Reference	C Assembly	C Assembly	GO Assembly
Implemented primitives	SIDH PKE KEM	SIDH KEM	SIDH KEM
Available parameters	p434 p503 p610 p751	p434 p503 p610 p751	p434 p503 p751
Optimized versions	Generic (portable c) x64 x64 compressed ARM64 ARM Cortex M4 VHDL	Generic (portable c) x64 x64 compressed ARMv8 ARMv8 compressed	Generic (GO) AMD64 ARM64
Security	Constant time	Constant time	Constant time

Table 2.1.: Overview and comparison of existing SIDH implementations.

3. Benchmarks

In this chapter the SIDH implementations introduced in chapter 2 are compared with each other. For this purpose, a benchmarking suite was developed, which allows the generation of comparable benchmarks between these implementations. To get a better understanding of the results, the chapter starts with section 3.1 describing the tools and methodology used for benchmarking. The following implementations details of section 3.2 provide precise internals of the benchmarking suite. This might be used for further development of the software. A description of how to actually use the presented benchmarking suite will be given in section 3.3. Finally, the benchmarking results are listed in section 3.4.

3.1. Benchmarking Methodology

In order to generate independent and stable benchmarking results, the benchmarking suite runs within a virtual environment: *Docker* is used to separate the running suite from the host operating system. This reduces the influence of resource intensive processes, which might falsify benchmarking results. Moreover, *Docker* enables a portable and scalable software solution.

In the following, the process of benchmarking is described within five steps. In a short version, the required steps are:

1. Create the benchmarking source code
2. Compile the benchmarking source code
3. Run *Callgrind*
4. Run *Massif*
5. Collect benchmarks

Create the benchmarking code

Each software libraries presented in chapter 2 implements various cryptographic primitives. To avoid overhead when calculating benchmarks, only the required functions to generate a SIDH key exchange should be called. Thus, a simple benchmarking file is provided for each implementation, which calls the appropriate underlying API. This benchmarking file must ensure that all required headers are imported, the API is called correctly and a main-function is provided.

Compile the benchmarking code

Once the benchmarking source code is created, it needs to be compiled to a binary. Therefore, all required dependencies (libraries, headers, source code, ..) need to be provided to the compiler. The benchmarking suite provides a Makefile for each implementation, where the compilation process is implemented. All compilations performed for the benchmarking suite must use consistent compiler optimizations. This ensures the comparability of the outputs. Currently, the optimization flag `-O3` is passed to the `gcc` compiler. Note, that CIRCL is implemented in GO and therefore the `go` compiler is used for compilation. Since this compiler does not provide optimization flags, the default compiler optimizations are used [34]. To be able to extract benchmarks for specific functions *inlining* is disabled during compilation. The C programming language offers the following directive to disable *inlining* for a function:

```
1 // This function will not be inlined by the compiler
2 void __attribute__((noinline)) no_inlining() {
3 // ...
4 }
5
6 // This function might be inlined by the compiler
7 void inlining() {
8 // ...
9 }
```

The `go` compiler can be directly invoked with a no-inlining (`-l`) flag:

```
1 go build -gcflags "-l"
```

Run Callgrind

Callgrind records function calls of a binary. For each call the executed instructions are counted. Moreover, the tool provides detailed information about the callee and how often functions are called. Thus, the tool also provides information about execution hotspots of the binary. *Callgrind* is part of *Valgrind*, a profiling tool which allows deep analysis of executed binaries. *Callgrind* is invoked on the command line via:

```
1 valgrind --tool=callgrind --callgrind-out-file=callgrind.out binary
```

The profiling data of *callgrind* is written to the file defined by `-callgrind-out-file`. This file might be analyzed using a graphical tool like *KCachegrind* or any other analyzing script. Running a binary using *callgrind*, slows down the execution times significantly. This is the main reason for the long execution times of the benchmarking suite.

Run Massif

Massif measures memory usage of a binary, including heap and stack. *Massif* is also part of the profiling tool *Valgrind*. The tool creates multiple snapshots of the memory consumption

during execution. Thus, one can extract the maximum memory consumption of a binary. The following command runs the *Massif*:

```
1 valgrind --tool=massif --stacks=yes --massif-out-file=massif.out binary
```

The profiling data will be written to the file defined by `-massif-out-file`. *Massif-visualizer* could be used to graphically analyze the data.

Collect benchmarks

Once the output files of *Callgrind* and *Massif* are produced, one can analyze the corresponding files to obtain:

1. Absolute instructions per function.
2. Maximum memory consumption during SIDH key exchange.

This information is finally used by the benchmarking suite, to produce graphs and tables for further investigation. To receive reliable information, all benchmarks are executed multiple times.

To further increase the quality and reproducibility of the results, the cache of the virtual operating system is cleared, before *Callgrind* and *Massif* are executed. This is done via:

```
1 sync; sudo sh -c "echo_1_>/proc/sys/vm/drop_caches"
2 sync; sudo sh -c "echo_2_>/proc/sys/vm/drop_caches"
3 sync; sudo sh -c "echo_3_>/proc/sys/vm/drop_caches"
```

3.2. Application Details

The benchmarking suite is developed in Python3 on a Linux/Ubuntu operating system. Currently, the following implementations are included:

- SIKE working on: p434, p503, p610, p751
 - Sike reference implementation (SIKE_Reference)
 - Sike generic optimized implementation (SIKE_Generic)
 - Sike generic optimized and compressed implementation (SIKE_Generic_Compressed)
 - Sike x64 optimized implementation (SIKE_x64)
 - Sike x64 optimized and compressed implementation (SIKE_x64_Compressed)
- PQCrypto-SIDH working on: p434, p503, p610, p751
 - PQCrypto-SIDH generic optimized implementation (Microsoft_Generic)
 - PQCrypto-SIDH generic optimized and compressed implementation (Microsoft_Generic_Compressed)
 - PQCrypto-SIDH x64 optimized implementation (Microsoft_x64)

- PQCrypto-SIDH x64 optimized and compressed implementation (Microsoft_x64_Compressed)
- CIRCL working on: p434, p503, p610, p751
 - CIRCL x64 optimized implementation (CIRCL_x64)

Furthermore, the benchmarking suite provides classical elliptic curve Diffie-Hellman (ECDH) based on OpenSSL. Since SIDH is a candidate to replace current Diffie-Hellman algorithms, ECDH is intended as reference value: The objective is it to compare optimized modern ECDH with quantum-secure SIDH.

Since each parameter set of SIDH meet a different security level (compare section 1.2.3) ECDH is also instantiated with different curves, each matching a SIDH security level. The used elliptic curves are namely:

1. secp256r1 (openssl: prime256v1 [35]): 128 bit security matching SIKEp434 [36]
2. secp384r1 (openssl: secp384r1): 192 bit security matching SIKEp610 [36]
3. secp521r1 (openssl: secp521r1): 256 bit security matching SIKEp751 [36]

For each of these introduced implementations the application measures benchmarks. In the following sections detailed information about the internals of the suite is given. Beside the precise application flow (subsection 3.2.1), the internal class structure of the Python3 code is shown (subsection 3.2.2).

3.2.1. Application Flow

This sections illustrates the application flow of the benchmarking suite (see Figure 3.1). Once triggered to run benchmarks, the following procedure is repeated for every implementation: The suite first compiles the benchmarking code. The binary is then executed multiple times to generate N benchmarking results, respectively for *Callgrind* and *Massif*.

Finally, the results are visualized in different formats. All resulting values are averages over N samples. Graphs compare the recorded instruction counts and the peak memory consumption among implementations instantiated with comparable security classes. The HTML and Latex tables lists all benchmarks per implementation and shows the standard derivation over N samples.

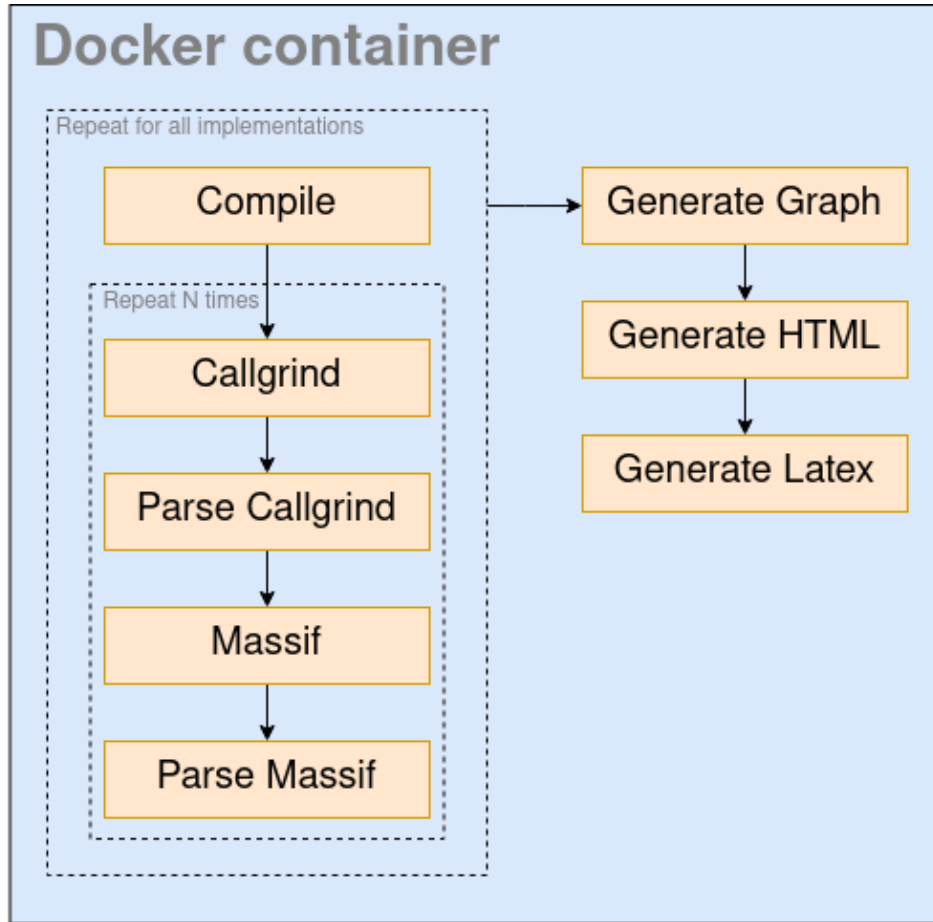


Figure 3.1.: Flow chart of the benchmarking suite. For each SIDH implementation, the source code is compiled and benchmarked multiple times. The benchmarking results are visualized in different format: Graphs, HTML and Latex.

3.2.2. Application Structure

This section covers the internal structure of the benchmarking suite. It illustrates, how each implementation is represented in code and how the benchmarking results are managed. To get a detailed description of the implemented functions have a look into the well-documented source code.

Representation of concrete implementations

The main logic of the benchmarking suite is placed within the class `BaseImplementation`. This class implements the logic of compiling code and measuring benchmarks. For each implementation which shall be benchmarked, a subclass of `BaseImplementation` is created (see Figure 3.2). These subclasses provide a link to the respective *Makefile*, which is used by the `BaseImplementation` to compile code, run `callgrind` and run `massif`. Furthermore, each

subclass can provide a set of arguments passed to the *Makefile*.

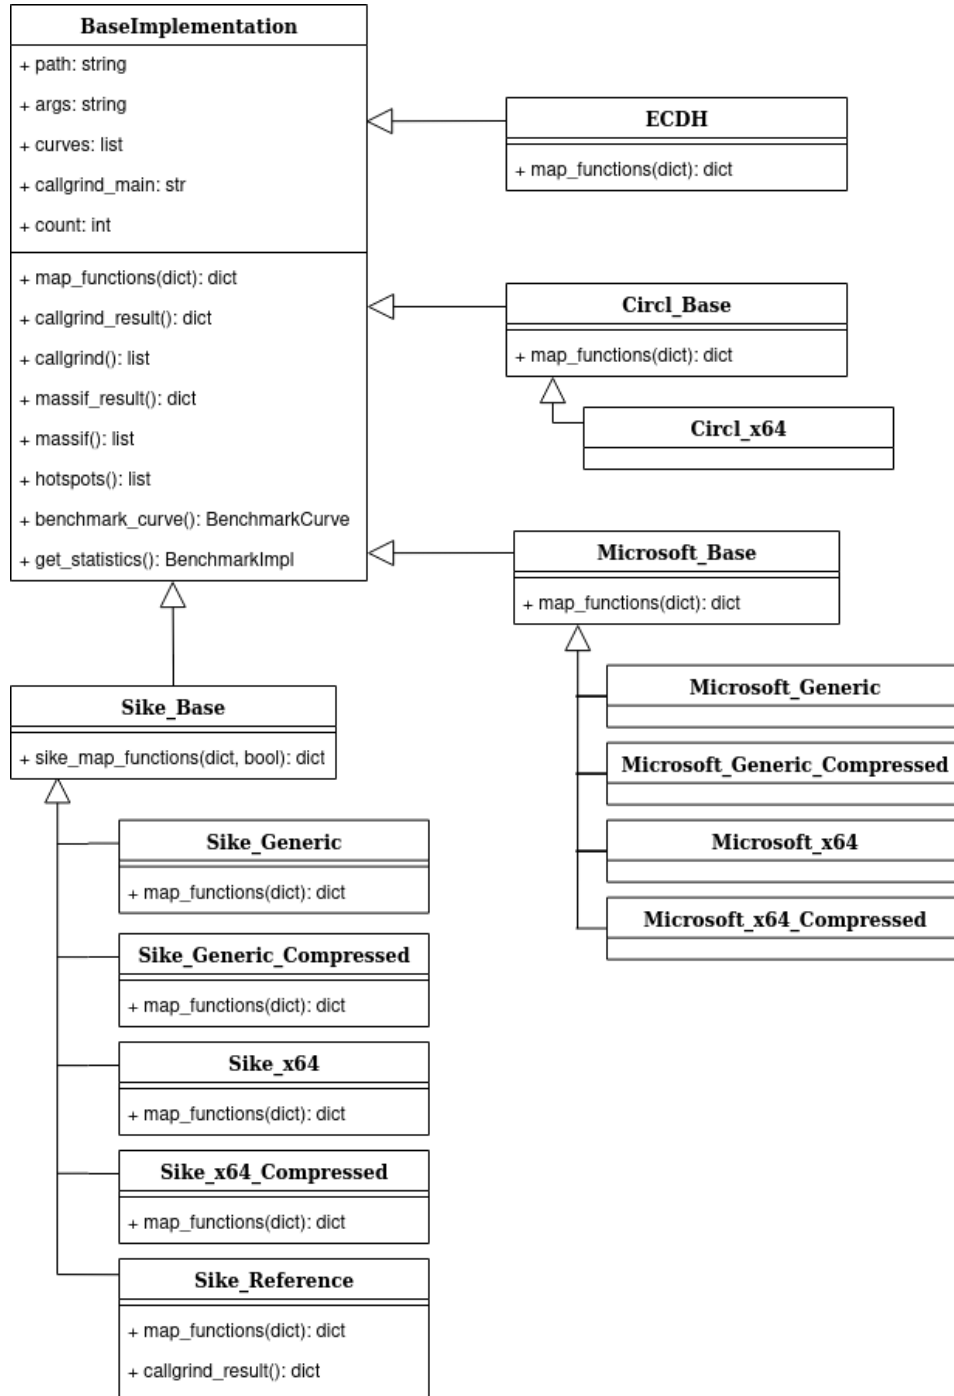


Figure 3.2.: Class diagram for the all implementations supported by the benchmarking suite. All concrete implementations of SIDH inherit from BaseImplementation. A *Makefile* provided by each subclass specifies how the benchmarking code is built and run.

Representation of benchmarking results

In order to ensure a clear analysis of the results the application implements a structure for the returned benchmarking values. Since each implementation can be run based on different parameter sets (in the terminology of the benchmarking suite this is called *curves*) and for each curve different benchmarking values are processed, the following hierarchy is applied (see Figure 3.3).

The class `BenchmarkImpl` represents the benchmarking results of a specific implementation. Therefore, the class manages a list of `BenchmarkCurve` objects, which contains benchmarks for a specific curve. Each benchmark for such a curve is represented by an instance of `Benchmark`. Thus, `BenchmarkCurve` holds a list of `Benchmark` objects.

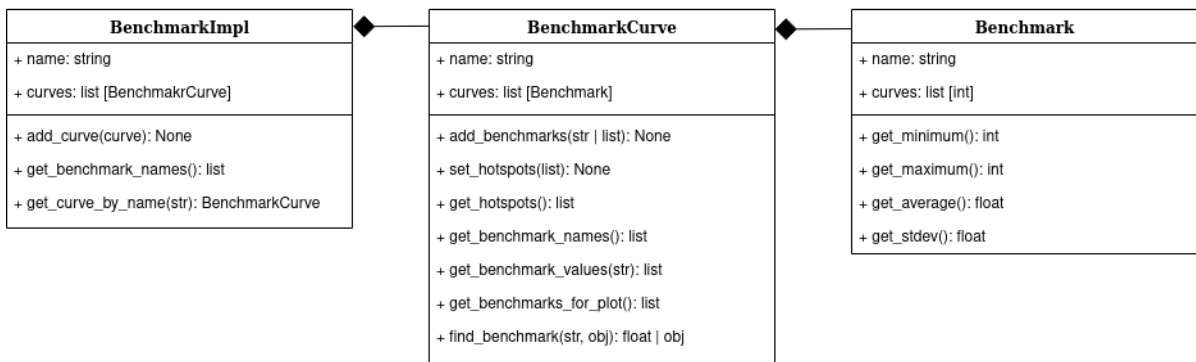


Figure 3.3.: Class diagram representing the management of the benchmarking results.

3.2.3. Adding Implementations

To add a new implementation to the benchmarking suite it is helpful to have a look at existing implementations. Adding a SIDH implementation to the benchmarking suite requires the following steps:

1. Add a new folder into the container/ folder of the root directory
container/new_implementation/.
2. Create a file container/new_implementation/benchmark.c that calls the SIDH key exchange API of the new implementation.
3. Add all necessary dependencies into container/new_implementation/ that are needed to compile benchmark.c. Note, maybe some changes in the *Dockerfile* are necessary to install certain dependencies on the virtual operating system started by Docker.
4. Create a Makefile, that supports the following commands:
 - build: Compile benchmark.c into the executable
new_implementation/build/benchmark
 - callgrind: Runs callgrind with the executable and stores the result in
new_implementation/benchmark/callgrind.out.

- massif: Runs massif with the executable and stores the result in `new_implementation/benchmark/massif.out`.

5. Add a new class to the source code, which inherits from `BaseImplementation`. You need to overwrite the `map_functions()` method to map the specific API functions of the new implementation to the naming used within the benchmarking suite. The new class might look similar to this:

```
1 class New_Implementation(BaseImplementation):
2     def __init__(self, count):
3         super().__init__(count=count,
4                           path=[path to Makefile],
5                           args=[args for Makefile],
6                           callgrind_main=[name of main function],
7                           curves=curves)
8
9     def map_functions(self, callgrind_result: dict) -> dict:
10        res = {
11            "PrivateKeyA": callgrind_result[[name of api function]],
12            "PublicKeyA": callgrind_result[[name of api function]],
13            "PrivateKeyB": callgrind_result[[name of api function]],
14            "PublicKeyB": callgrind_result[[name of api function]],
15            "SecretA": callgrind_result[[name of api function]],
16            "SecretB": callgrind_result[[name of api function]],
17        }
18        return res
```

6. Import the created class into `container/benchmarking.py` and add the class to the `implementations` list:

```
1 implementations =[
2     #...
3     New_Implementation,
4 ]
```

Once these steps are done the benchmarking suite is able to benchmark the new implementation.

3.3. Usage

This section explains the usage of the benchmarking suite. Beside the execution of benchmarks, the application also provides unit tests for the implementation. Both tasks, running benchmarks and executing unit tests, need to run within the docker container. Thus, it is

mandatory to install docker on your system to run the benchmarking suite¹ To provide a easy to use interface for both tasks a script *run.sh* is available. This script can be run with different arguments:

- Building the docker container:

```
./run.sh build
```

- Running unit tests within the docker container.

```
./run.sh test
```

Testing the benchmarking suite runs for a while, since each implementation is compiled to verify the functionality of the *Makefiles*. However, this procedure is logged while the unit tests are executed.

- Running benchmarks within the docker container.

```
./run.sh benchmark
```

This command triggers the *benchmarking.py* script within the docker container. This script contains the entry logic of the benchmarking suite. It benchmarks all available implementations and generates different output formats (graphs, html, latex). It is configurable how often each benchmark should be measured: The variable *N* in *container/benchmarking.py* describes the amount of repetitions for *callgrind* and *massif*. Once the benchmarks are done, all output files can be inspected in the folder *data/* of your current directory. These files visualize average values over *N* samples. The output files are:

- *XXX.png*: These files compare the absolute instruction count of all implementations, which were run on the *XXX* parameter set ($XXX \in \{434, 503, 610, 751\}$). Additionally, it contains a ECDH benchmark for comparison.
- *XXX_mem.png*: These files compare the peak memory consumption of all implementations, which were run on the *XXX* parameter set ($XXX \in \{434, 503, 610, 751\}$). Additionally, it contains a ECDH benchmark for comparison.
- *cached.json*: This cache file contains all benchmarking results. It can be used as input for the benchmarking suite to use cached data instead of generating all benchmarks again. Since the benchmarking suite runs multiple hours if each implementation is evaluated *N*=100 times, this functionality provides great speed up if only the output data should change. To use the cached file as input, copy it to *container/.cached/cached.json* and run the benchmarking suite again.
- *results.html*: This file lists detailed benchmarking results in a human readable HTML table.

¹Instructions for installing *Docker*: <https://docs.docker.com/get-docker/>
(The application was developed using Docker version 19.03.8, build afacb8b7f0)

- *results.tex*: This file contains the benchmarking results formatted in a latex table. The output is used for this document.

3.4. Results

The results presented in this section were calculated on a x64 architecture (Intel(R) Core(TM) i5-6200U CPU @ 2.30GHz) running Ubuntu 20.04.1 LTS. The installed docker version was 19.03.8. The benchmarking suite was initialized to run *callgrind* and *massif* 100 times, respectively (e.g. $N=100$). This section visualizes the results of the benchmarking process. The docker container was invoked with the following command:

```
./run.sh benchmark
```

The exact values measured (averages and standard derivation over $N=100$ samples) and the identified execution hotspots for each implementation described in section 3.2 can be found in addenda A.1.

This section visualizes the performance of existing SIDH libraries (*SIKE*, *PQ-Crypto (Microsoft)* and *CIRCL*). Note, that this section does not interpret any results. A detailed analysis is given in chapter 4. When looking at the benchmarks of all the different implementations, one need to consider the characteristics of *optimized* (generic and x64) implementations, *compressed* implementations and *reference* implementations. The *reference* implementation by SIKE is by far the slowest version of the benchmarked SIDH key exchanges. *Optimized* versions exploit hardware specific instructions leading to significantly reduced instruction counts. All *compressed* versions benefit from smaller key sizes, however, they suffer from increased execution times and intense memory allocations.

To get a better feeling of the speed of compressed versions and the reference implementation compared to optimized versions, all SIKE implementations are shown in subsection 3.4.1 (initialized with p434, however the chosen parameter set is not relevant for this purpose). Since the current SIDH libraries are also compared to the highly optimized ECDH openssl library, the figures in subsection 3.4.2 only consider the optimized implementations of the libraries. Finally, subsection 3.4.3 compares all compressed implementations of the SIDH key exchange.

All graphs in this section compare the average instruction count and the peak memory consumption among implementations instantiated with comparable security classes. Since the difference between implementations does marginally change if considering different security classes and to avoid many graphs this section only the highest and lowest security classes are visualized:

- NIST security level 1 (~AES 128 bit) represented by p434 and specp256r1
- NIST security level 5 (~AES 256 bit) represented by p751 and specp521r1.

3.4.1. SIKE Implementations

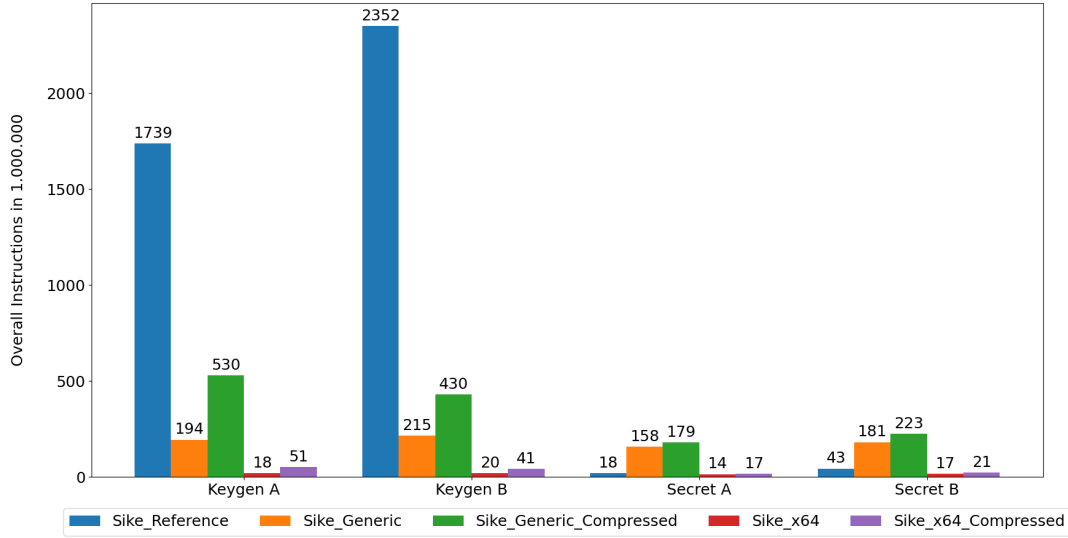


Figure 3.4.: Overall instructions for all SIKE implementations initialized with p434. The reference implementation is the slowest, the x64 optimized version is the fastest. These results meet intuitive expectations.

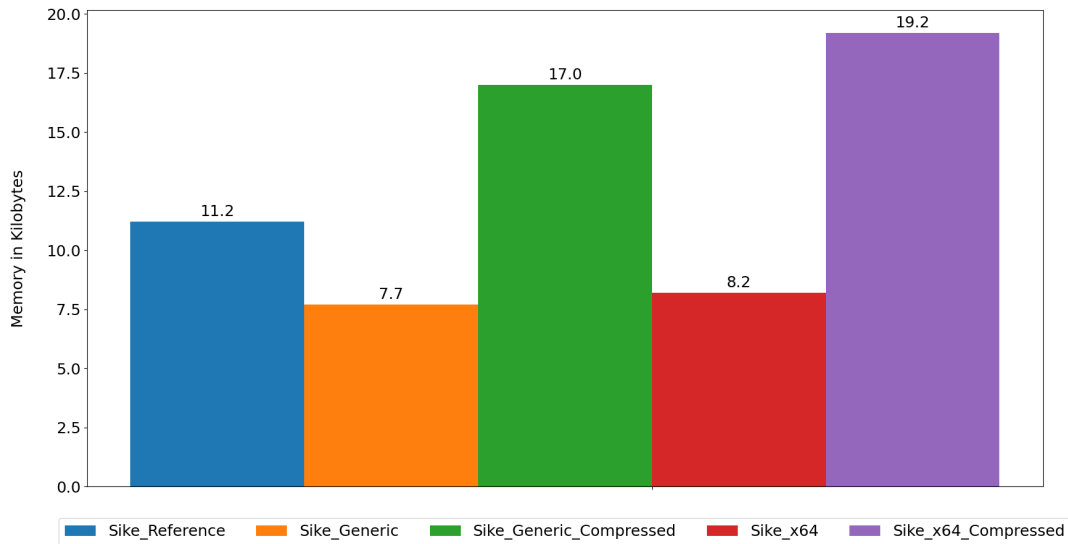


Figure 3.5.: Maximum memory consumption in kilobytes for all SIKE implementations initialized with p434. The required memory overhead of compressed versions is clearly visible.

3.4.2. Optimized Implementations

Results matching NIST security level 1

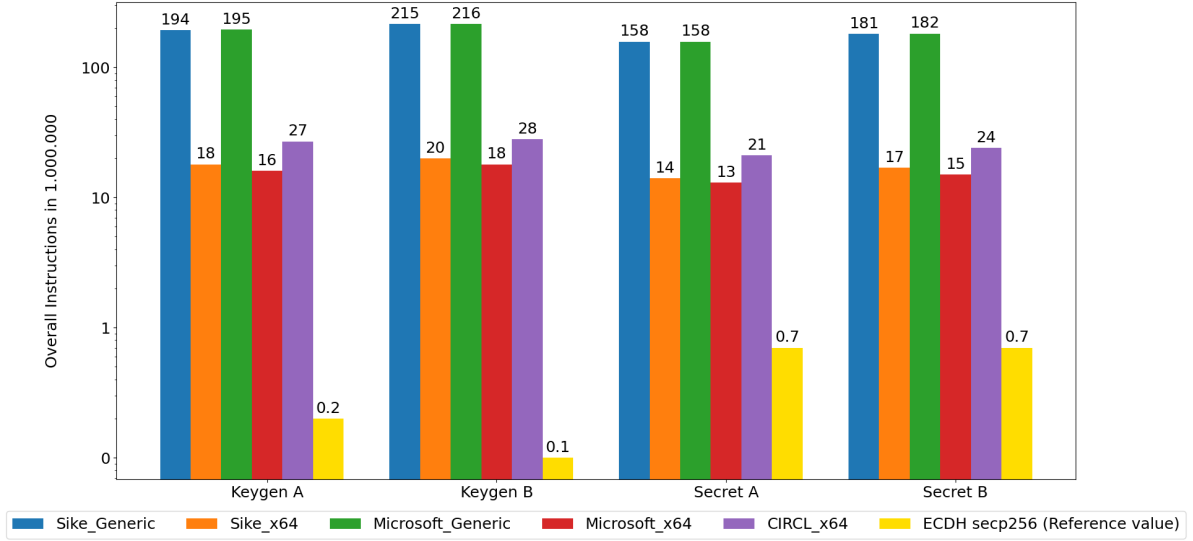


Figure 3.6.: Overall instructions for SIDH parameter p434 compared to ECDH via secp256r1.

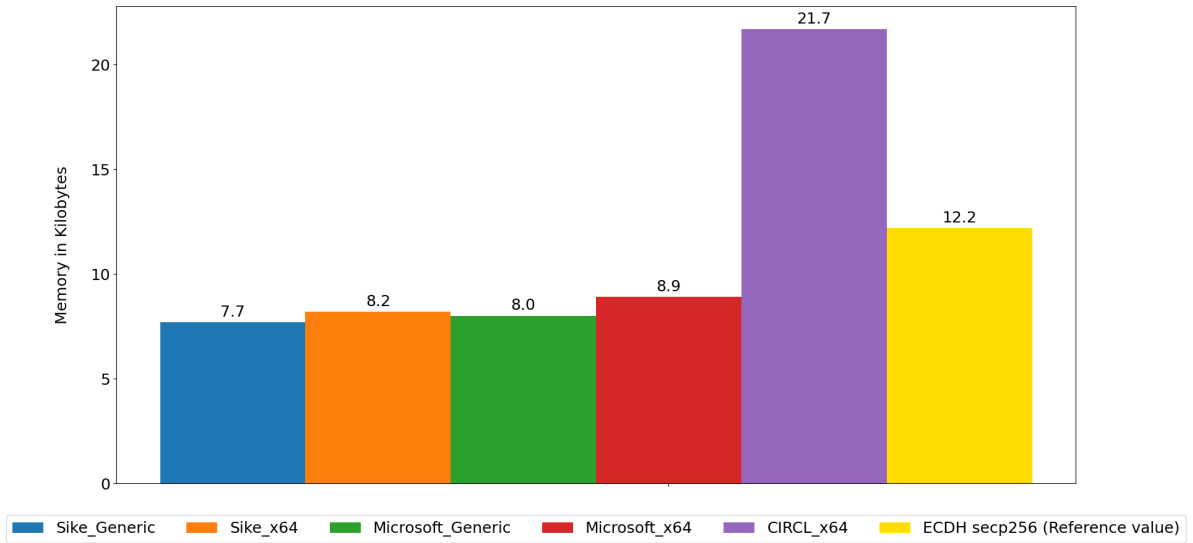


Figure 3.7.: Maximum memory consumption in kilobytes for SIDH parameter p434 compared to ECDH via secp256r1.

Results matching NIST security level 5

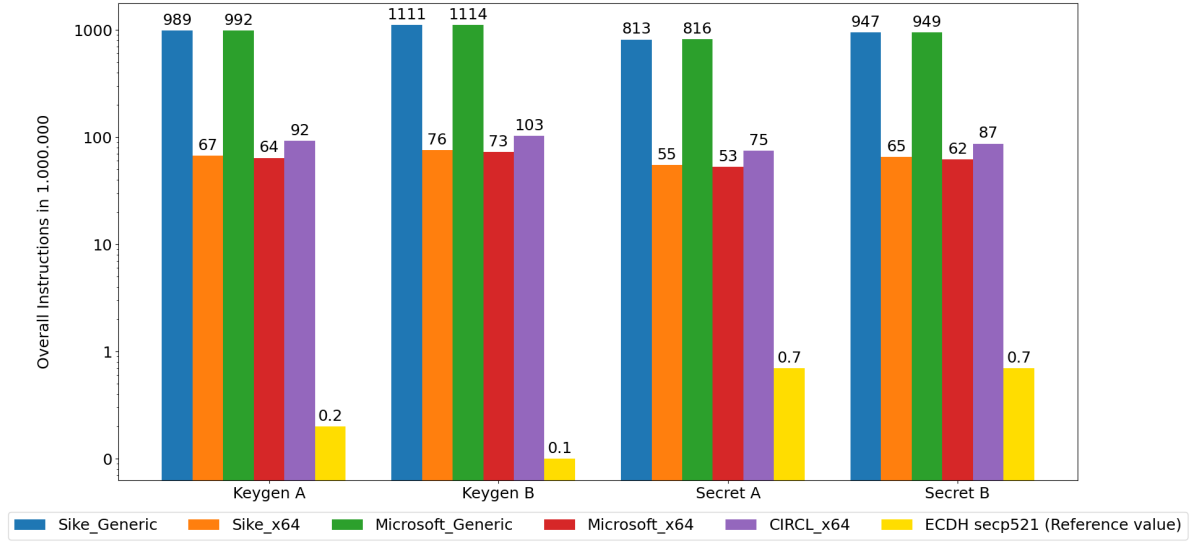


Figure 3.8.: Overall instructions for SIDH parameter p751 compared to ECDH via secp521r1.

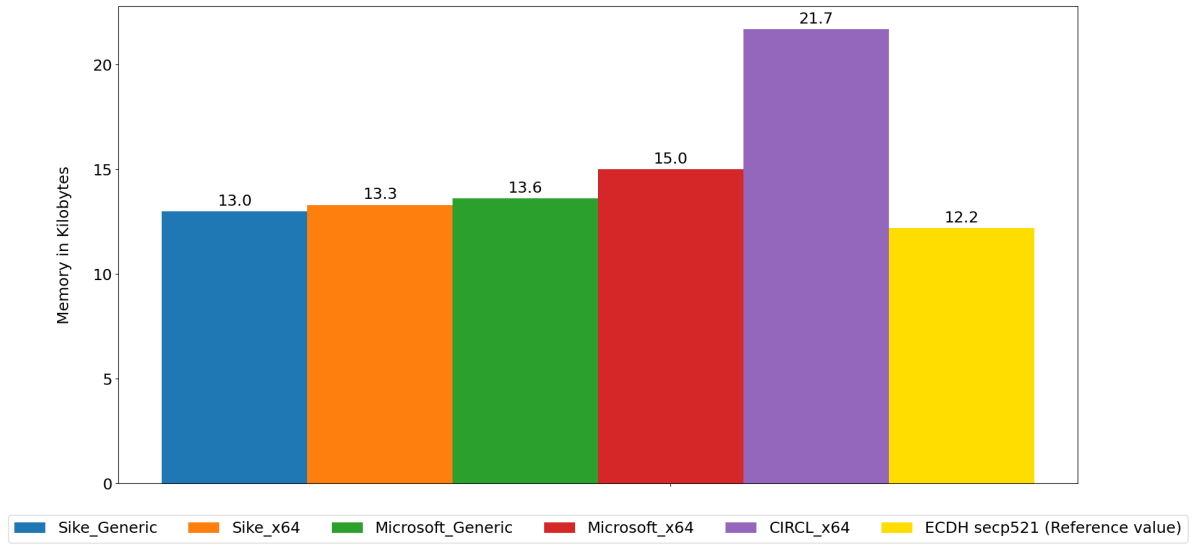


Figure 3.9.: Maximum memory consumption in kilobytes for SIDH parameter p751 compared to ECDH via secp521r1.

3.4.3. Compressed implementations

Results matching NIST security level 1

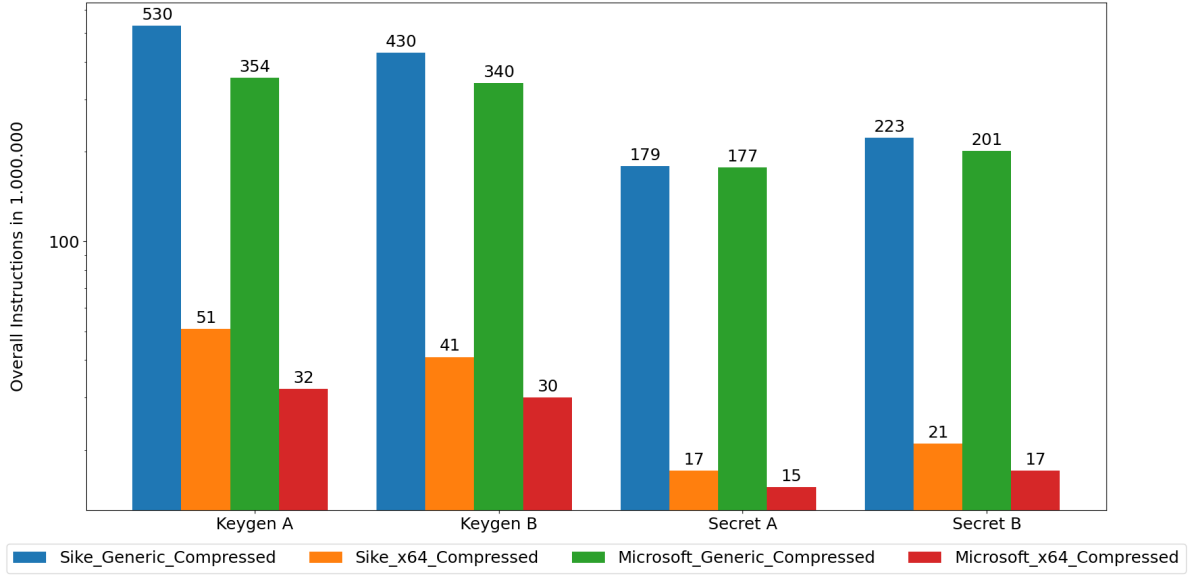


Figure 3.10.: Overall instructions for compressed SIDH parameter p434.

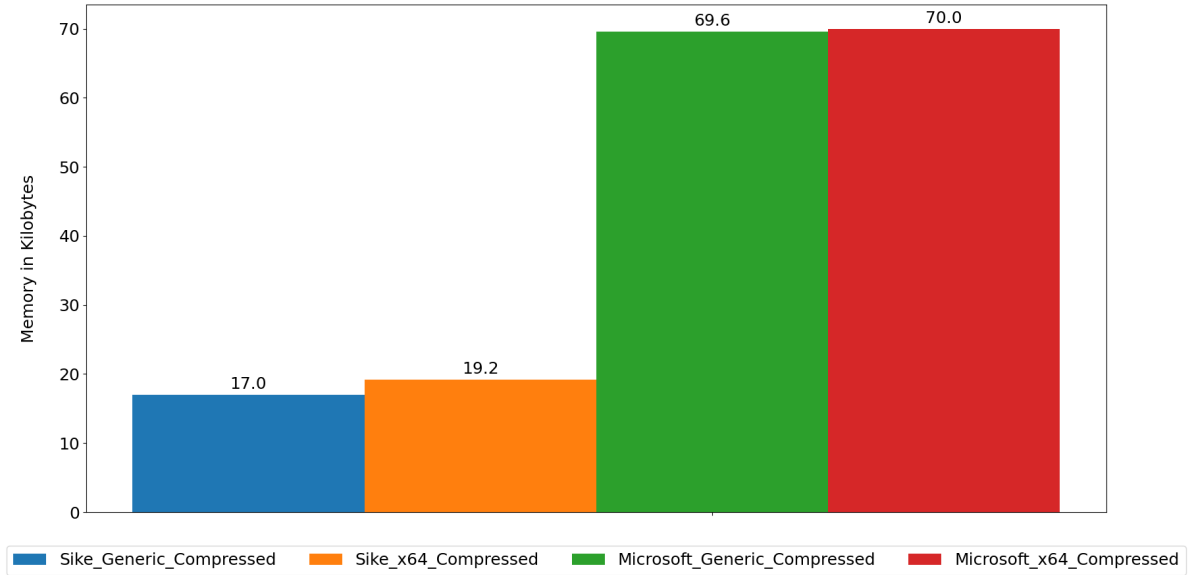


Figure 3.11.: Maximum memory consumption in kilobytes for compressed SIDH parameter p434.

Results matching NIST security level 5

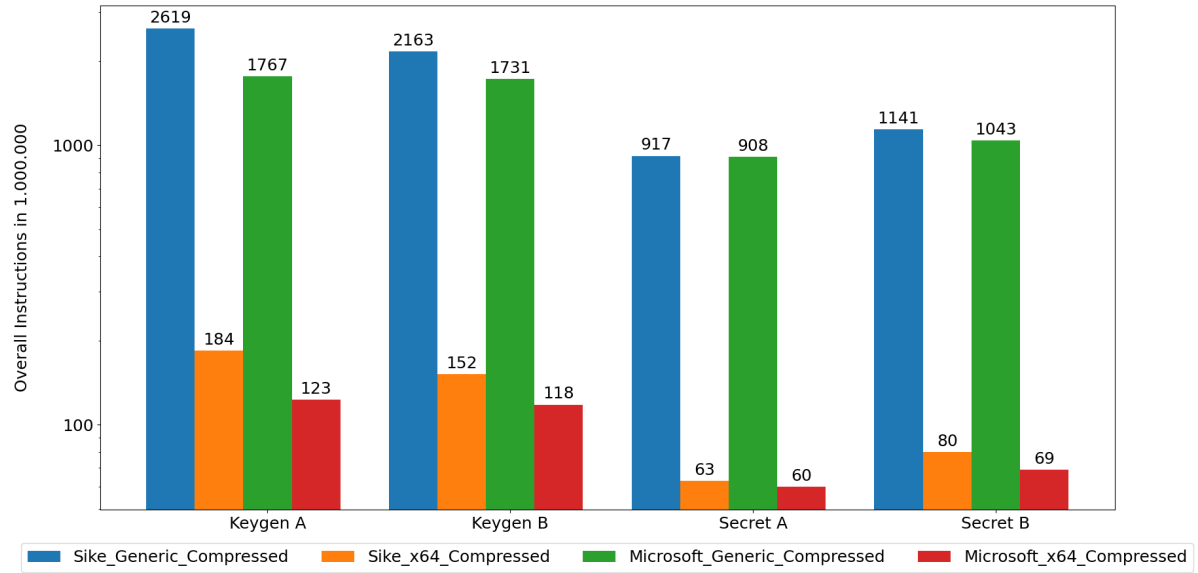


Figure 3.12.: Overall instructions for compressed SIDH parameter p751.

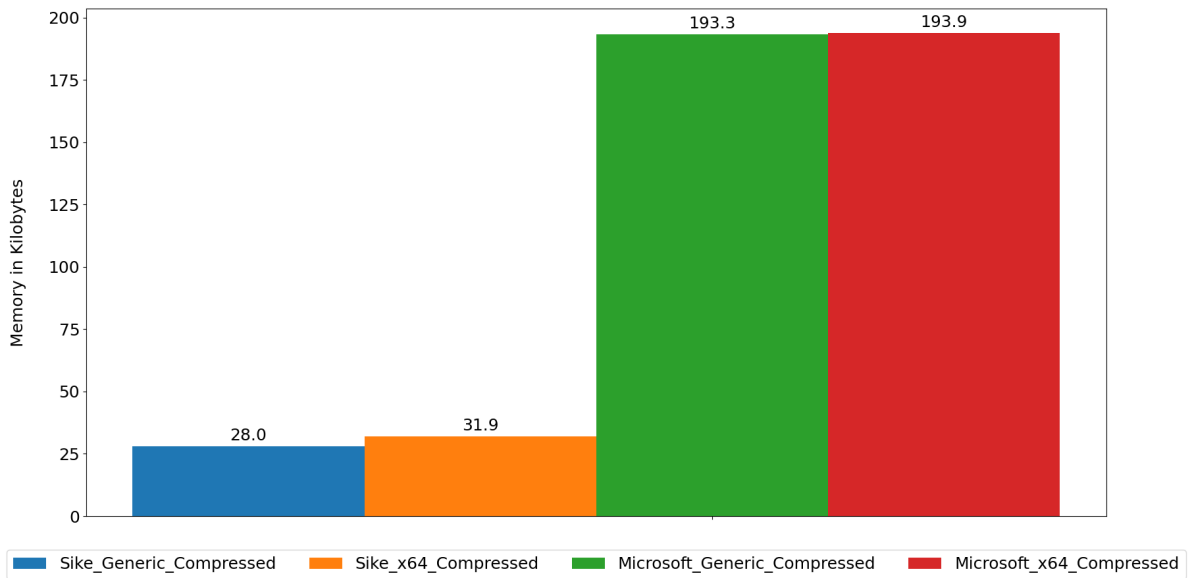


Figure 3.13.: Maximum memory consumption in kilobytes for compressed SIDH parameter p751.

4. Performance Analysis

This chapter analyses the results presented in section 3.4. First of all the efficiency in terms of memory consumption and execution runtime is investigated (section 4.1). The security of the benchmarked libraries will be evaluated in section 4.2.

4.1. Measures for Efficiency

The efficiency of SIDH libraries will be quantified by the peak memory consumption and by the absolute instruction count. These measures are benchmarked for all implemented parameter sets of a given implementation (for details see chapter 3). subsection 4.1.1 firstly compares SIDH libraries among each other before modern key exchange primitives (namely ECDH) are taken into account in subsection 4.1.2.

4.1.1. Comparing SIDH libraries

Before comparisons between the SIDH libraries *SIKE*, *PQ-Crypto* and *CIRCL* are drawn, the performance differences between *optimized* and *compressed* versions are investigated. Exemplarily, all variants of SIKE are taken into considerations. Note that all *compressed* versions of SIKE also contain *optimized* code. Their key sizes, however, are reduced compared to non-compressed variants.

Figure 3.4 clearly shows a difference between *optimized* and *compressed* versions of SIKE in term of absolute instructions. Compressed versions roughly claim twice as much operations (~500 million for *SIKE_Generic_Compressed* and ~45 million *SIKE_x64_Compressed*) to generate a key pair as non-compressed variants (~200 million for *SIKE_Generic* and ~19 million *SIKE_x64*). Additionally, the generation of the shared secret is a little slower for the compressed implementations.

Besides execution times, Figure 3.5 also compares the memory consumption. As stated by the authors of SIKE, compressed variants allocate more memory: The peak memory allocation is with 17 kilobytes (*SIKE_Generic_Compressed*) and 19.2 kilobytes (*SIKE_x64_Compressed*) twice as high as the non-compressed versions.

The *SIKE_Reference* implementation is with 11.2 kB allocated memory peak close to the average of all SIKE implementations. However, regarding the execution times for key generation, the reference implementation is by far the slowest variant: The more than 2.3 billion (2.300.000.000) instructions were measured in order to generate Bobs key pair. This is slower by factor 100 compared to *SIKE_x64*. As the name suggests, this reference implementation must be seen as a proof-of-concept. Thus, it will not be considered further in this analysis.

To sum up, the comparison between all SIKE implementations showed, that *optimized* versions have decreased instructions counts as well as decreased memory consumptions compared to *compressed* implementations. The further analysis of the detailed benchmarks in in addenda A.1 lead to a very similar result for the *PQ-Crypto* library of Microsoft. One can state, that *compressed* versions roughly demand twice as much resources than *optimized* variants.

Generic optimized implementations

Generic optimized implementations contain generally optimized code for various hardware platforms. However, no hardware specific instructions can be exploited in these versions. The benchmarking suite currently supports ¹:

1. *SIKE_Generic*
2. *Microsoft_Generic*

Figure 3.6 shows the benchmarks for execution time of all optimized variants initiated with p_{434} . For this section, however, only *SIKE_Generic* and *Microsoft_Generic* are taken into account. In all four steps of the SIDH key exchange the benchmarked values for both implementations look almost the same. However, when considering exact measurements from A.1, one can observe that *SIKE_Generic* executes constantly half a million operations less than *Microsoft_Generic* (this holds for all four categories: Keygen A, Keygen B, Secret A and Secret B). While this sounds significant, the relative difference is actually less than 0.01%. While the absolute gap further increases if higher security classes are analyzed (about three million operations constant difference for p_{751}), the relative disparity stays smaller than 0.01%.

Peak memory consumptions for parameter p_{434} are visualized in Figure 3.7. As in terms of execution time, memory allocation numbers between *SIKE_Generic* and *Microsoft_Generic* hardly differ: The SIKE version occupies 0.3 kB less memory for p_{434} and 0.7 kB less than for p_{751} . Overall, the relative difference regarding memory consumption is about 5%.

Although both implementations hardly differ in their performance benchmarks one can state, that *SIKE_Generic* demands less resources than *Microsoft_Generic*. This disparity, however, is marginal.

x64 optimized implementations

X64 optimized implementations exploit AMD64 specific hardware operations to improve performance on these machines. The benchmarking suite currently supports the following x64 optimized variants:

1. *SIKE_x64*
2. *Microsoft_x64*

¹Although CIRCL states to offer a generic optimized implementation[32], this version could not be compiled.

3. CIRCL_x64

Figure 3.6 shows the execution time benchmarks for all optimized variants initiated with $p434$. For this section, however, only *SIKE_x64*, *Microsoft_x64* and *CIRCL_x64* are taken into account. In all four categories listed in the graph *Microsoft_x64* is the fastest implementation. *SIKE_x64* is slightly slower executing about two million instructions more in each category. This corresponds to a relative difference of about 10%. The most expensive implementation in terms of performed operations is *CIRCL_x64*: 40% more operations are needed in each step of the SIDH key exchange compared to the fastest variant *Microsoft_x64*.

Figure 3.8 compares the implementations initiated with $p751$ - matching the highest NIST security level 5. *Microsoft_x64* stays the fastest version while the relative difference to *SIKE_x64* (5%) and *CIRCL_x64* (30%) decreases.

To compare the memory consumption of the x64 optimized implementations, have a look at Figure 3.7. The memory benchmarks of *SIKE_x64* (8.2 kB) and *Microsoft_x64* (8.9 kB) hardly differ, whereas *CIRCL_x64* has a peak allocation of 24.7 kB for a single SIDH key exchange. This is by factor 2.5 greater compared to the others. However, Figure 3.9 reveals that the memory consumption for *CIRCL_x64* does barely change for higher security classes. Nevertheless, *CIRCL_x64* has the most intense memory consumption of all x64 optimized implementations and *Microsoft_x64* allocates roughly about 10% more memory than *SIKE_x64* (this holds respectively for all security classes).

The fastest x64 optimized SIDH key exchange is performed by the Microsoft library *PQ-Crypto*. *SIKE_x64* is slightly slower but allocates less memory than *Microsoft_x64*. The most resources are consumed by *CIRCL_x64* (instruction count and memory allocations).

Analysis of execution hotspots

4.1.2. Comparing SIDH and ECDH

Overall

Analysis of ECDH execution hotspots

4.2. Security Considerations

In order to analyze the given implementations in terms of security, the claim of all libraries to implement security relevant functions in constant time is investigated in subsection 4.2.1. The implemented key sizes of all SIDH libraries will be considered additionally in subsection 4.2.2.

4.2.1. Constant time

4.2.2. Key length

5. Conclusion

6. Introduction

Use with pdfLaTeX and Biber.

6.1. Section

Citation test (with Biber) [latex].

6.1.1. Subsection

See Table 6.1, Figure 6.1, Figure 6.2, Figure 6.3, Figure 6.4, Figure 6.5.

Table 6.1.: An example for a simple table.

A	B	C	D
1	2	1	2
2	3	2	3

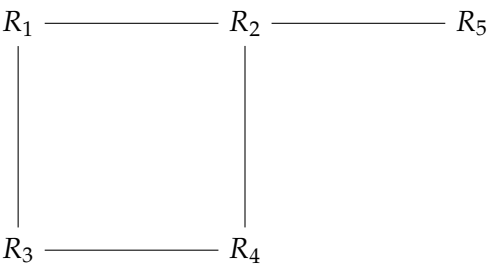


Figure 6.1.: An example for a simple drawing.

This is how the glossary will be used.

Donor dye, ex. Alexa 488 (D_{dye}), Förster distance, Förster distance (R_0), and k_{DEAC} . Also, the TUM has many computers, not only one Computer. Subsequent acronym usage will only print the short version of Technical University of Munich (TUM) (take care of plural, if needed!), like here with TUM, too. It can also be \rightarrow hidden¹ $<-$.

[(TODO: Now it is your turn to write your thesis.

This will be a few tough weeks.)]

[(DONE: NEVERTHELESS, CELEBRATE IT WHEN IT IS DONE!)]

¹Example for a hidden TUM glossary entry.

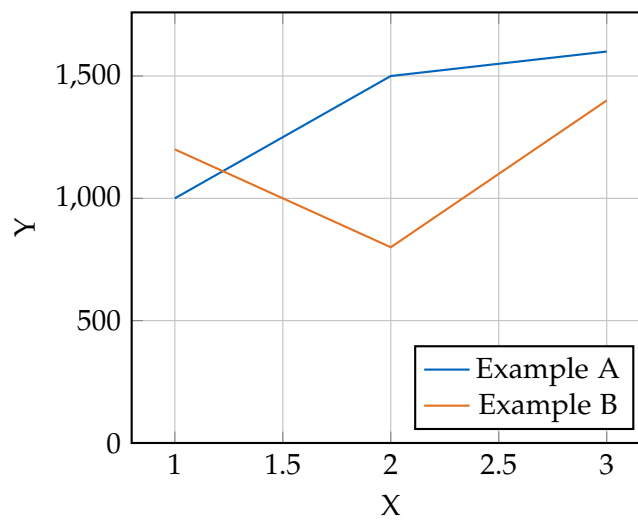


Figure 6.2.: An example for a simple plot.

```
1 SELECT * FROM tbl WHERE tbl.str = "str"
```

Figure 6.3.: An example for a source code listing.



Figure 6.4.: Includegraphics searches for the filename without extension first in logos, then in figures.



Figure 6.5.: For pictures with the same name, the direct folder needs to be chosen.



(a) The logo.



(b) The famous slide.

Figure 6.6.: Two TUM pictures side by side.

A. General Addenda

A.1. Detailed Benchmarks

A.1.1. Benchmarks for ECDH

Parameter	secp256	secp384	secp521
PrivateKeyA	0 (0)	0 (0)	0 (0)
PublicKeyA	159.418 (0)	159.418 (0)	159.418 (0)
PrivateKeyB	0 (0)	0 (0)	0 (0)
PublicKeyB	114.430 (0)	114.430 (0)	114.430 (1)
SecretA	652.796 (0)	652.796 (0)	652.796 (0)
SecretB	651.270 (0)	651.270 (0)	651.270 (0)
Memory in bytes	12.152 (0)	12.152 (0)	12.152 (0)

Table A.1.: Benchmarks for ECDH

Execution hotspots parameter *secp256*:

1. `__ecp_nistz256_mul_montq`: 29.89%
2. `__ecp_nistz256_sqr_montq`: 18.32%
3. `_dl_relocate_object`: 7.27%

Execution hotspots parameter *secp384*:

1. `__ecp_nistz256_mul_montq`: 29.89%
2. `__ecp_nistz256_sqr_montq`: 18.32%
3. `_dl_relocate_object`: 7.27%

Execution hotspots parameter *secp521*:

1. `__ecp_nistz256_mul_montq`: 29.89%
2. `__ecp_nistz256_sqr_montq`: 18.32%
3. `_dl_relocate_object`: 7.27%

A.1.2. Benchmarks for Sike Reference

Parameter	434	503	610	751
PrivateKeyA	28.919 (0)	29.020 (0)	34.541 (0)	34.770 (0)
PublicKeyA	1.739.736.057 (0)	2.512.464.770 (0)	4.308.288.592 (0)	7.461.795.045 (0)
PrivateKeyB	29.418 (0)	29.519 (0)	34.617 (0)	35.289 (0)
PublicKeyB	2.352.757.723 (0)	3.435.254.160 (0)	5.790.924.796 (0)	10.556.125.964 (0)
SecretA	18.726.146 (0)	22.075.791 (0)	27.927.622 (0)	35.617.111 (0)
SecretB	43.530.260 (0)	56.573.382 (0)	79.580.027 (0)	118.687.930 (0)
Memory in bytes	11.208 (0)	11.928 (0)	12.904 (0)	13.736 (0)

Table A.2.: Benchmarks for Sike Reference

Execution hotspots parameter 434:

1. `__gmpn_sbpi1_div_qr`: 10.7%
2. `__gmpn_tdiv_qr`: 9.59%
3. `__gmpn_hgcd2`: 9.29%

Execution hotspots parameter 503:

1. `__gmpn_sbpi1_div_qr`: 11.08%
2. `__gmpn_hgcd2`: 9.8%
3. `__gmpn_submul_1`: 9.61%

Execution hotspots parameter 610:

1. `__gmpn_submul_1`: 12.08%
2. `__gmpn_sbpi1_div_qr`: 11.67%
3. `__gmpn_mul_basecase`: 10.13%

Execution hotspots parameter 751:

1. `__gmpn_submul_1`: 15.21%
2. `__gmpn_mul_basecase`: 11.82%
3. `__gmpn_sbpi1_div_qr`: 11.69%

A.1.3. Benchmarks for Sike Generic

Parameter	434	503	610	751
PrivateKeyA	90 (0)	95 (0)	96 (0)	97 (0)
PublicKeyA	194.932.002 (0)	299.462.858 (0)	618.958.459 (0)	989.778.051 (0)
PrivateKeyB	57 (0)	57 (0)	53 (0)	59 (0)
PublicKeyB	215.702.626 (0)	329.835.556 (0)	616.667.764 (0)	1.111.885.426 (0)
SecretA	158.061.535 (0)	243.950.880 (0)	515.975.033 (0)	813.862.458 (0)
SecretB	181.708.340 (0)	278.459.825 (0)	522.486.909 (0)	947.216.296 (0)
Memory in bytes	7.720 (0)	7.784 (0)	11.288 (0)	13.016 (0)

Table A.3.: Benchmarks for Sike Generic

Execution hotspots parameter 434:

1. mp_mul : 59.45%
2. rdc_mont : 31.67%
3. fp2mul434_mont : 4.58%

Execution hotspots parameter 503:

1. mp_mul : 59.06%
2. rdc_mont : 33.02%
3. fp2mul503_mont : 4.07%

Execution hotspots parameter 610:

1. mp_mul : 60.05%
2. rdc_mont : 32.8%
3. fp2mul610_mont : 4.0%

Execution hotspots parameter 751:

1. mp_mul : 61.06%
2. rdc_mont : 32.76%
3. fp2mul751_mont : 3.42%

A.1.4. Benchmarks for Sike Generic Compressed

Parameter	434	503	610	751
PrivateKeyA	97 (0)	95 (0)	99 (0)	107 (0)
PublicKeyA	530.693.055 (36.636.736)	800.568.613 (51.179.821)	1.546.560.342 (104.069.626)	2.619.265.805 (178.208.196)
PrivateKeyB	185 (0)	145 (0)	215 (0)	200 (0)
PublicKeyB	430.783.226 (2.956.687)	648.550.068 (5.818.409)	1.209.676.275 (9.989.150)	2.163.046.938 (23.915.892)
SecretA	179.835.016 (2.751)	276.744.609 (2.820)	575.396.837 (4.432)	917.385.785 (6.137)
SecretB	223.183.731 (2.871.219)	342.669.712 (3.247.052)	640.870.333 (5.431.240)	1.141.966.654 (14.089.895)
Memory in bytes	16.968 (0)	19.080 (0)	23.976 (0)	28.008 (0)

Table A.4.: Benchmarks for Sike Generic Compressed

Execution hotspots parameter 434:

1. mp_mul : 58.77%
2. rdc_mont : 32.15%
3. fp2mul434_mont : 4.13%

Execution hotspots parameter 503:

1. mp_mul : 58.4%
2. rdc_mont : 33.46%
3. fp2mul503_mont : 3.71%

Execution hotspots parameter 610:

1. mp_mul : 59.42%
2. rdc_mont : 33.34%
3. fp2mul610_mont : 3.61%

Execution hotspots parameter 751:

1. mp_mul : 60.46%
2. rdc_mont : 33.29%
3. fp2mul751_mont : 3.1%

A.1.5. Benchmarks for Sike x64

Parameter	434	503	610	751
PrivateKeyA	90 (0)	95 (0)	96 (0)	97 (0)
PublicKeyA	18.197.636 (0)	25.309.825 (0)	46.870.491 (0)	67.976.631 (0)
PrivateKeyB	57 (0)	57 (0)	53 (0)	59 (0)
PublicKeyB	20.227.975 (0)	28.024.313 (0)	46.946.162 (0)	76.798.386 (0)
SecretA	14.735.273 (0)	20.595.673 (0)	39.015.534 (0)	55.840.033 (0)
SecretB	17.044.799 (0)	23.672.739 (0)	39.800.369 (0)	65.465.094 (0)
Memory in bytes	8.168 (0)	8.520 (0)	11.392 (0)	13.328 (0)

Table A.5.: Benchmarks for Sike x64

Execution hotspots parameter 434:

1. mp_mul : 52.23%
2. rdc_mont : 23.46%
3. fpsub434 : 4.22%

Execution hotspots parameter 503:

1. mp_mul : 49.88%
2. rdc_mont : 27.18%
3. fpsub503 : 4.12%

Execution hotspots parameter 610:

1. mp_mul : 51.51%
2. rdc_mont : 28.57%
3. fpsub610 : 3.72%

Execution hotspots parameter 751:

1. mp_mul : 53.34%
2. rdc_mont : 27.94%
3. fpsub751 : 3.81%

A.1.6. Benchmarks for Sike x64 Compressed

Parameter	434	503	610	751
PrivateKeyA	97 (0)	95 (0)	99 (0)	107 (0)
PublicKeyA	51.478.796 (3.842.012)	70.633.473 (4.623.263)	120.698.795 (6.557.301)	184.983.750 (10.400.022)
PrivateKeyB	142 (0)	144 (0)	179 (0)	188 (0)
PublicKeyB	41.648.989 (345.077)	56.629.973 (517.293)	94.418.639 (796.824)	152.512.949 (1.142.155)
SecretA	17.062.490 (1.558)	23.745.084 (2.055)	44.067.290 (2.871)	63.697.941 (4.288)
SecretB	21.428.968 (241.953)	29.717.372 (262.099)	49.633.365 (394.707)	80.095.836 (886.898)
Memory in bytes	19.240 (0)	21.608 (0)	27.264 (0)	31.936 (0)

Table A.6.: Benchmarks for Sike x64 Compressed

Execution hotspots parameter 434:

1. mp_mul : 50.21%
2. rdc_mont : 23.14%
3. fpsub434 : 4.22%

Execution hotspots parameter 503:

1. mp_mul : 47.88%
2. rdc_mont : 26.74%
3. fpsub503 : 4.1%

Execution hotspots parameter 610:

1. mp_mul : 49.73%
2. rdc_mont : 28.31%
3. fpsub610 : 3.75%

Execution hotspots parameter 751:

1. mp_mul : 51.71%
2. rdc_mont : 27.79%
3. fpsub751 : 3.82%

A.1.7. Benchmarks for CIRCL x64

Parameter	434	503	751
PrivateKeyA	671 (0)	671 (0)	671 (0)
PublicKeyA	27.173.063 (12.785)	30.588.687 (14.339)	92.516.170 (11.132)
PrivateKeyB	672 (0)	672 (0)	672 (0)
PublicKeyB	28.959.178 (431)	32.771.303 (314)	103.101.351 (244)
SecretA	21.089.046 (2.101)	23.926.495 (1.461)	75.130.980 (636)
SecretB	24.364.188 (324)	27.627.692 (360)	87.853.968 (337)
Memory in bytes	21.654 (1.022)	21.691 (896)	21.663 (1.055)

Table A.7.: Benchmarks for CIRCL x64

Execution hotspots parameter 434:

1. p434.mulP434: 47.28%
2. p434.rdcP434: 22.93%
3. p434.subP434: 5.81%

Execution hotspots parameter 503:

1. p503.mulP503: 41.39%
2. p503.rdcP503: 23.04%
3. p503.subP503: 8.39%

Execution hotspots parameter 751:

1. p751.mulP751: 56.16%
2. p751.rdcP751: 21.72%
3. p751.subP751: 6.02%

A.1.8. Benchmarks for Microsoft Generic

Parameter	434	503	610	751
PrivateKeyA	86 (0)	91 (0)	91 (0)	91 (0)
PublicKeyA	195.425.484 (0)	300.437.100 (0)	620.117.514 (0)	992.618.213 (0)
PrivateKeyB	53 (0)	53 (0)	48 (0)	53 (0)
PublicKeyB	216.131.501 (0)	330.771.955 (0)	617.536.325 (0)	1.114.734.257 (0)
SecretA	158.459.759 (0)	244.758.599 (0)	516.977.970 (0)	816.142.131 (0)
SecretB	182.033.988 (0)	279.212.025 (0)	523.130.069 (0)	949.500.928 (0)
Memory in bytes	8.040 (0)	8.360 (0)	11.784 (0)	13.624 (0)

Table A.8.: Benchmarks for Microsoft Generic

Execution hotspots parameter 434:

1. mp_mul : 60.55%
2. rdc_mont : 31.9%
3. fp2mul434_mont : 4.56%

Execution hotspots parameter 503:

1. mp_mul : 60.12%
2. rdc_mont : 33.25%
3. fp2mul503_mont : 3.96%

Execution hotspots parameter 610:

1. mp_mul : 61.24%
2. rdc_mont : 32.54%
3. fp2mul610_mont : 4.02%

Execution hotspots parameter 751:

1. mp_mul : 62.22%
2. rdc_mont : 32.44%
3. fp2mul751_mont : 3.43%

A.1.9. Benchmarks for Microsoft Generic Compressed

Parameter	434	503	610	751
PrivateKeyA	97 (0)	95 (0)	99 (0)	105 (0)
PublicKeyA	354.807.131 (28.640.848)	548.929.246 (48.233.475)	1.058.952.592 (69.882.412)	1.767.280.284 (114.624.073)
PrivateKeyB	239 (0)	233 (0)	206 (0)	314 (0)
PublicKeyB	340.645.715 (4.510.751)	513.777.907 (5.790.443)	956.709.613 (11.260.973)	1.731.204.963 (18.353.464)
SecretA	177.456.197 (1.970)	274.002.754 (1.875)	569.195.853 (3.954)	908.459.651 (5.191)
SecretB	201.752.486 (871)	309.529.792 (1.000)	577.403.156 (1.459)	1.043.916.680 (1.618)
Memory in bytes	69.576 (0)	89.896 (0)	134.600 (0)	193.336 (0)

Table A.9.: Benchmarks for Microsoft Generic Compressed

Execution hotspots parameter 434:

1. mp_mul : 59.82%
2. rdc_mont : 32.48%
3. fp2mul434_mont : 4.06%

Execution hotspots parameter 503:

1. mp_mul : 59.26%
2. rdc_mont : 33.89%
3. fp2mul503_mont : 2.89%

Execution hotspots parameter 610:

1. mp_mul : 60.56%
2. rdc_mont : 33.26%
3. fp2mul610_mont : 3.55%

Execution hotspots parameter 751:

1. mp_mul : 61.48%
2. rdc_mont : 33.03%
3. fp2mul751_mont : 2.33%

A.1.10. Benchmarks for Microsoft x64

Parameter	434	503	610	751
PrivateKeyA	86 (0)	91 (0)	91 (0)	91 (0)
PublicKeyA	16.733.523 (0)	23.373.795 (0)	44.826.467 (0)	64.976.610 (0)
PrivateKeyB	53 (0)	53 (0)	48 (0)	53 (0)
PublicKeyB	18.587.638 (0)	25.858.758 (0)	44.876.850 (0)	73.377.700 (0)
SecretA	13.529.422 (0)	18.993.109 (0)	37.346.394 (0)	53.326.381 (0)
SecretB	15.655.268 (0)	21.831.732 (0)	38.025.823 (0)	62.514.039 (0)
Memory in bytes	8.936 (0)	9.048 (0)	12.944 (0)	15.008 (0)

Table A.10.: Benchmarks for Microsoft x64

Execution hotspots parameter 434:

1. mp_mul : 55.81%
2. rdc_mont : 23.82%
3. 0x000000000000cd3c : 4.56%

Execution hotspots parameter 503:

1. mp_mul : 52.33%
2. rdc_mont : 28.27%
3. 0x000000000000d8e4 : 4.38%

Execution hotspots parameter 610:

1. mp_mul : 53.86%
2. rdc_mont : 29.88%
3. 0x000000000000f256 : 3.86%

Execution hotspots parameter 751:

1. mp_mul : 55.83%
2. rdc_mont : 29.24%
3. 0x000000000000fefc : 3.62%

A.1.11. Benchmarks for Microsoft x64 Compressed

Parameter	434	503	610	751
PrivateKeyA	97 (0)	95 (0)	99 (0)	105 (0)
PublicKeyA	32.099.997 (2.309.282)	44.418.194 (3.712.967)	80.996.244 (6.287.779)	123.157.832 (12.599.138)
PrivateKeyB	149 (0)	152 (0)	187 (0)	200 (0)
PublicKeyB	30.784.457 (316.772)	41.848.170 (319.398)	72.151.350 (669.781)	118.144.325 (1.757.118)
SecretA	15.538.409 (588)	21.689.738 (483)	42.096.902 (1.350)	60.202.926 (1.955)
SecretB	17.949.468 (571)	24.856.517 (656)	42.885.495 (895)	69.857.944 (1.243)
Memory in bytes	69.960 (0)	90.640 (0)	135.216 (0)	193.872 (0)

Table A.11.: Benchmarks for Microsoft x64 Compressed

Execution hotspots parameter 434:

1. mp_mul : 52.82%
2. rdc_mont : 23.28%
3. 0x000000000000247dc : 3.56%

Execution hotspots parameter 503:

1. mp_mul : 49.75%
2. rdc_mont : 27.75%
3. 0x00000000000026694 : 3.45%

Execution hotspots parameter 610:

1. mp_mul : 51.53%
2. rdc_mont : 29.48%
3. 0x0000000000002e046 : 3.09%

Execution hotspots parameter 751:

1. mp_mul : 53.4%
2. rdc_mont : 28.99%
3. 0x00000000000032f6d : 2.8%

List of Figures

1.1. Symmetric encryption scheme	1
1.2. Asymmetric encryption scheme	2
1.3. Diffie Hellman diagram	4
1.4. Supersingular Isogeny Diffie Hellman diagram	11
1.5. SIDH based on <i>isogen</i> and <i>isoex</i>	12
1.6. SIKE public key encryption	13
1.7. SIKE public key encryption	14
3.1. Flow chart of the benchmarking suite.	23
3.2. Class diagram for the all implementations supported by the benchmarking suite.	25
3.3. Class diagram for benchmarking results.	26
3.4. Overall instructions SIKE	30
3.5. Maximum memory consumption SIKE	30
3.6. Overall instructions p434	31
3.7. Maximum memory consumption p434	31
3.8. Overall instructions p751	32
3.9. Maximum memory consumption 751	32
3.10. Overall instructions compressed p434	33
3.11. Maximum memory consumption compressed p434	33
3.12. Overall instructions compressed p751	34
3.13. Maximum memory consumption compressed p751	34
6.1. Example drawing	39
6.2. Example plot	40
6.3. Example listing	40
6.4. Something else can be written here for listing this, otherwise the caption will be written!	40
6.5. For pictures with the same name, the direct folder needs to be chosen.	41
6.6. Two TUM pictures side by side.	41

List of Tables

1.1. Impact of quantum computers on modern encryption schemes	7
1.2. Core functions of the SIKE reference implementation	12
2.1. Existing SIDH implementations	18
6.1. Example table	39
A.1. Benchmarks for ECDH	42
A.2. Benchmarks for Sike Reference	43
A.3. Benchmarks for Sike Generic	44
A.4. Benchmarks for Sike Generic Compressed	45
A.5. Benchmarks for Sike x64	46
A.6. Benchmarks for Sike x64 Compressed	47
A.7. Benchmarks for CIRCL x64	48
A.8. Benchmarks for Microsoft Generic	49
A.9. Benchmarks for Microsoft Generic Compressed	50
A.10. Benchmarks for Microsoft x64	51
A.11. Benchmarks for Microsoft x64 Compressed	52

Bibliography

- [1] C. Eckert. *IT-Sicherheit*. Berlin, Boston: De Gruyter Oldenbourg, 21 Aug. 2018. ISBN: 978-3-11-056390-0. DOI: <https://doi.org/10.1515/9783110563900>. URL: <https://www.degruyter.com/view/title/530046>.
- [2] W. Diffie and M. Hellman. "New directions in cryptography". In: *IEEE transactions on Information Theory* 22.6 (1976), pp. 644–654.
- [3] D. Wätjen. *Kryptographie*. Springer, 2018.
- [4] J. Randall, B. Kaliski, J. Brainard, and S. Turner. "Use of the RSA-KEM Key Transport Algorithm in the Cryptographic Message Syntax (CMS)". In: *Proposed Standard 5990* (2010).
- [5] M. A. Nielsen and I. Chuang. *Quantum computation and quantum information*. 2002.
- [6] L. Chen, L. Chen, S. Jordan, Y.-K. Liu, D. Moody, R. Peralta, R. Perlner, and D. Smith-Tone. *Report on post-quantum cryptography*. Vol. 12. US Department of Commerce, National Institute of Standards and Technology, 2016.
- [7] D. A. Lidar and T. A. Brun. *Quantum error correction*. Cambridge university press, 2013.
- [8] F. Boudot, P. Gaudry, A. Guillevis, N. Heninger, E. Thome, and P. Zimmermann. *795-bit factoring and discrete logarithms*.
- [9] A. Beutelspacher, H. B. Neumann, and T. Schwarzpaul. "Der diskrete Logarithmus, Diffie-Hellman-Schlüsselvereinbarung, ElGamal-Systeme". In: *Kryptografie in Theorie und Praxis*. Springer, 2010, pp. 132–144.
- [10] P. W. Shor. "Algorithms for quantum computation: discrete logarithms and factoring". In: *Proceedings 35th annual symposium on foundations of computer science*. Ieee. 1994, pp. 124–134.
- [11] L. K. Grover. "A fast quantum mechanical algorithm for database search". In: *Proceedings of the twenty-eighth annual ACM symposium on Theory of computing*. 1996, pp. 212–219.
- [12] V. Mavroedis, K. Vishi, M. D. Zych, and A. Jøsang. "The impact of quantum computing on present cryptography". In: *arXiv preprint arXiv:1804.00200* (2018).
- [13] D. P. Chi, J. W. Choi, J. San Kim, and T. Kim. "Lattice based cryptography for beginners." In: *IACR Cryptol. ePrint Arch.* 2015 (2015), p. 938.
- [14] J. Hoffstein, J. Pipher, and J. H. Silverman. "NTRU: A ring-based public key cryptosystem". In: *International Algorithmic Number Theory Symposium*. Springer. 1998, pp. 267–288.

- [15] C. Gentry and D. Boneh. *A fully homomorphic encryption scheme*. Vol. 20. 9. Stanford university Stanford, 2009.
- [16] J. Hartmanis. “Computers and intractability: a guide to the theory of NP-completeness (michael r. garey and david s. johnson)”. In: *Siam Review* 24.1 (1982), p. 90.
- [17] J. Ding and A. Petzoldt. “Current state of multivariate cryptography”. In: *IEEE Security & Privacy* 15.4 (2017), pp. 28–36.
- [18] J. Ding and D. Schmidt. “Rainbow, a new multivariable polynomial signature scheme”. In: *International Conference on Applied Cryptography and Network Security*. Springer. 2005, pp. 164–175.
- [19] D. J. Bernstein and T. Lange. “Post-quantum cryptography”. In: *Nature* 549.7671 (2017), pp. 188–194.
- [20] R. J. McEliece. “A public-key cryptosystem based on algebraic”. In: *Coding Thv* 4244 (1978), pp. 114–116.
- [21] G. Becker. “Merkle signature schemes, merkle trees and their cryptanalysis”. In: *Ruhr-University Bochum, Tech. Rep* (2008).
- [22] R. C. Merkle. *Secrecy, authentication, and public key systems*. Stanford University, 1979.
- [23] D. Jao and L. De Feo. “Towards quantum-resistant cryptosystems from supersingular elliptic curve isogenies”. In: *International Workshop on Post-Quantum Cryptography*. Springer. 2011, pp. 19–34.
- [24] L. De Feo. “Supersingular Isogeny Key Encapsulation”. In: *NIST round 3 submission*. 2020.
- [25] D. Urbanik. “A Friendly Introduction to Supersingular Isogeny Diffie-Hellman”. In: (2017).
- [26] C. Costello. “Supersingular Isogeny Key Exchange for Beginners”. In: *International Conference on Selected Areas in Cryptography*. Springer. 2019, pp. 21–50.
- [27] C. Costello and C. AIMSCS. “A gentle introduction to isogeny-based cryptography”. In: *Tutorial Talk at SPACE* (2016).
- [28] S. Jaques and J. M. Schanck. “Quantum cryptanalysis in the RAM model: Claw-finding attacks on SIKE”. In: *Annual International Cryptology Conference*. Springer. 2019, pp. 32–61.
- [29] S. Tani. “Claw finding algorithms using quantum walk”. In: *Theoretical Computer Science* 410.50 (2009), pp. 5285–5297.
- [30] P. C. Van Oorschot and M. J. Wiener. “Parallel collision search with cryptanalytic applications”. In: *Journal of cryptology* 12.1 (1999), pp. 1–28.
- [31] Microsoft. *PQCrypto-SIDH*. <https://github.com/microsoft/PQCrypto-SIDH>. 2020.
- [32] Cloudflare. *CIRCL*. <https://github.com/cloudflare/circl>. 2020.
- [33] K. Kwiatkowski and A. Faz-Hernández. *Introducing CIRCL: An Advanced Cryptographic Library*. July 2019. URL: <https://blog.cloudflare.com/introducing-circl/>.

- [34] *GoLang Wiki Compiler Optimizations*. <https://github.com/golang/go/wiki/CompilerOptimizations>. Accessed: 2020-09-25.
- [35] S. Turner, D. Brown, K. Yiu, R. Housley, and T. Polk. "Elliptic curve cryptography subject public key information". In: *RFC 5480 (Proposed Standard)* (2009).
- [36] D. R. Brown. "Sec 2: Recommended elliptic curve domain parameters". In: *Standards for Efficient Cryptography* (2010).